

# ArcaneFEM: Scalable GPGPU Implicit FEM Solver for Unstructured Meshes

**Mohd Afeef Badri<sup>1,✉</sup>, G. Gospellier<sup>2</sup>, C.A. Leger<sup>1</sup>, K. Alaire<sup>1</sup>, E. Foerster<sup>1</sup>, A. l'Heritier<sup>2</sup>**

<sup>1</sup> CEA DES/ISAS/DM2S/SGLS, SACLAY, FRANCE

<sup>2</sup> CEA DAM, BRUYÈRES-LE-CHÂTEL, FRANCE

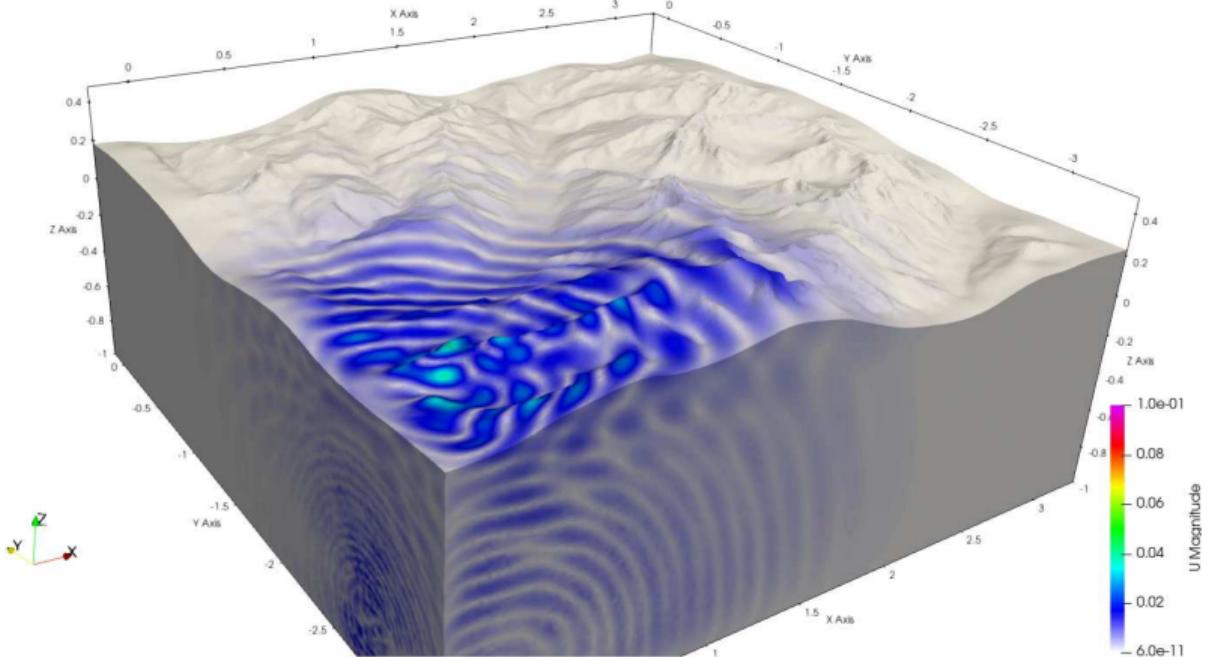
✉ MOHD-AFEEF.BADRI@CEA.FR



# ArcaneFEM: introduction

- Arcane + FEM started Nov. 2022 (**Young code**)
- Primary aim, **high-fidelity Earthquake** (source to site)
- ML Assisted Data-Assimilation: Fast solving

4M DOF Chamonix French Alps Simulation on 8 MPI in 22 min

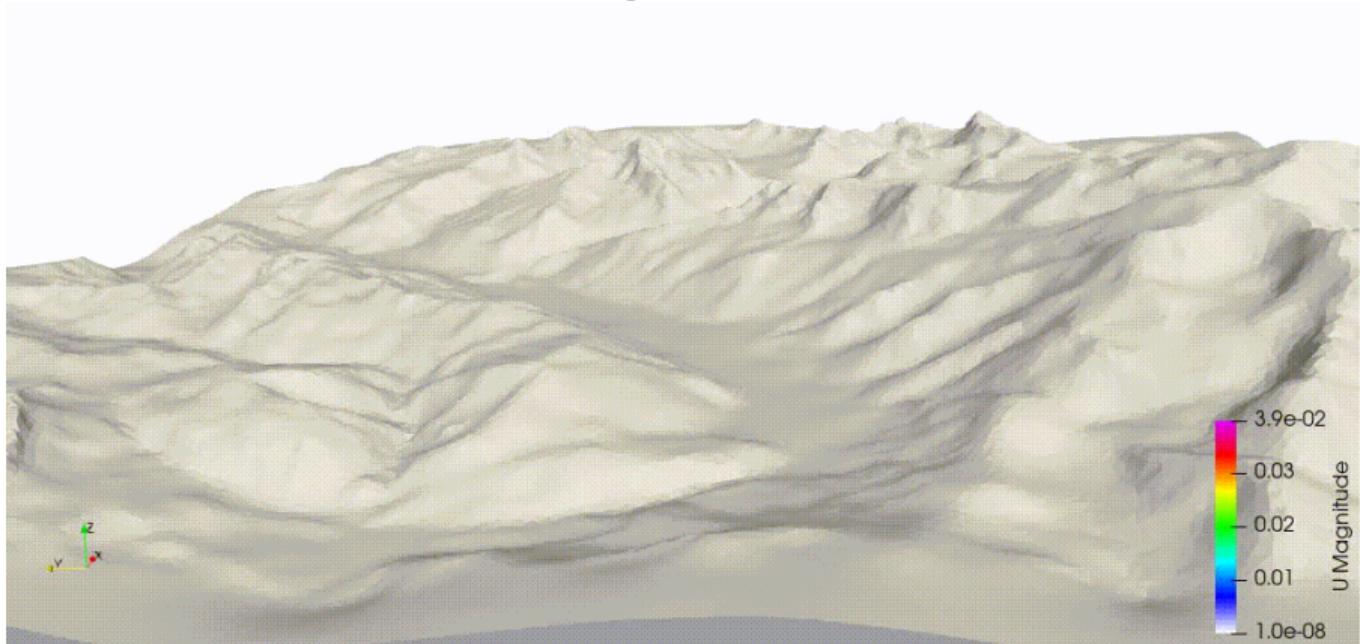




# ArcaneFEM: introduction

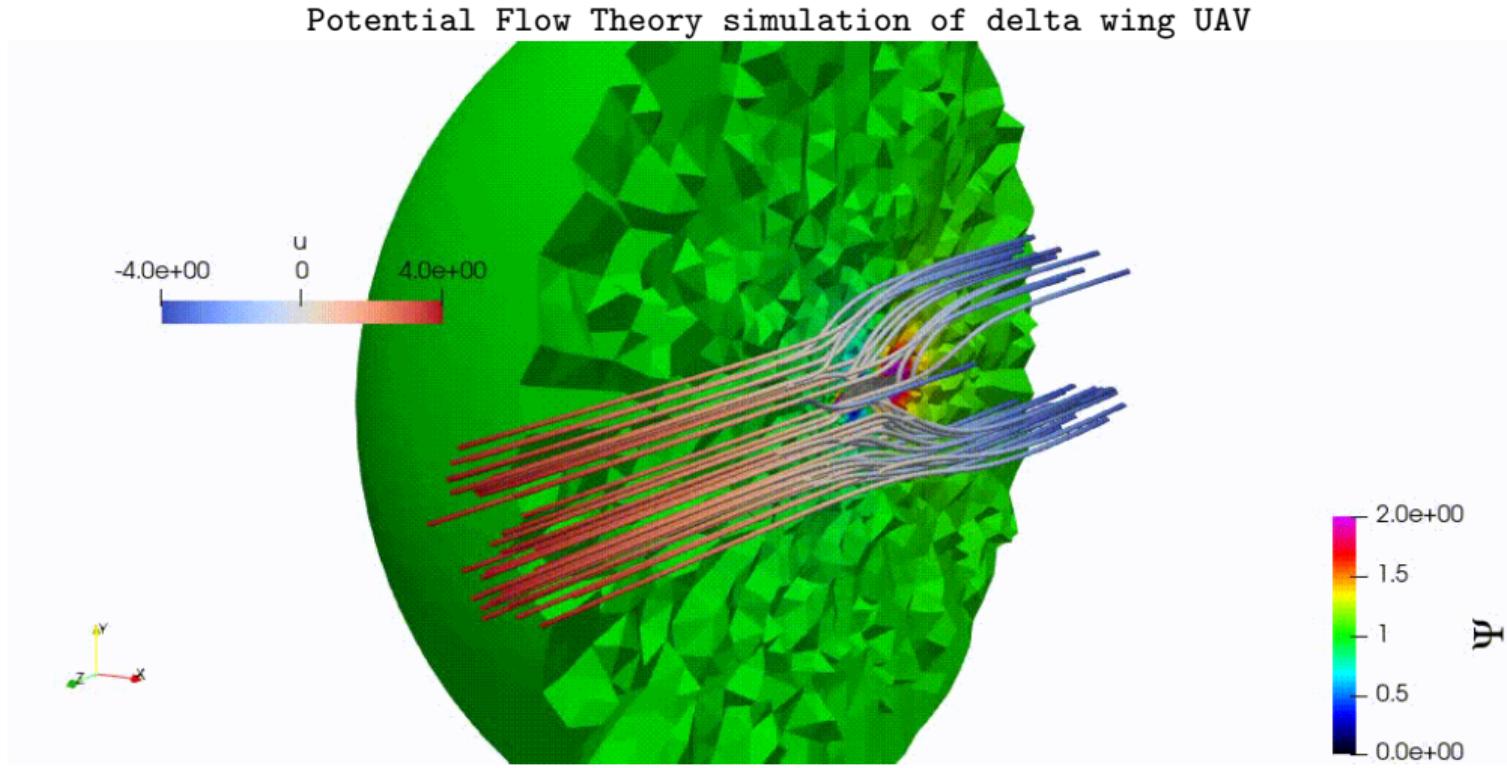
- Arcane + FEM started Nov. 2022 (**Young code**)
- Primary aim, **high-fidelity Earthquake** (source to site)
- ML Assisted Data-Assimilation: Fast solving

4M DOF Chamonix French Alps Simulation on 8 MPI in 22 min



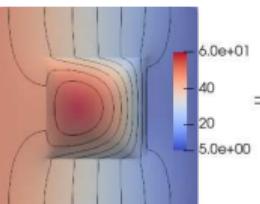
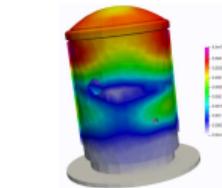
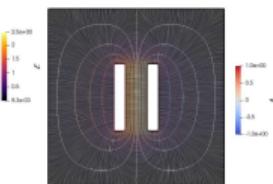
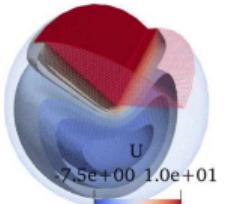
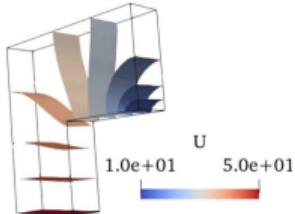
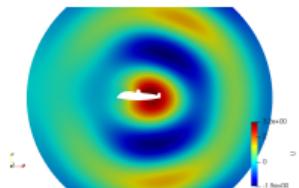
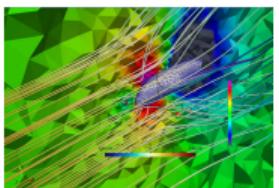
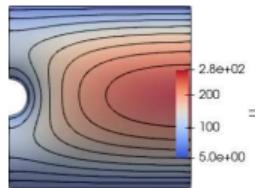
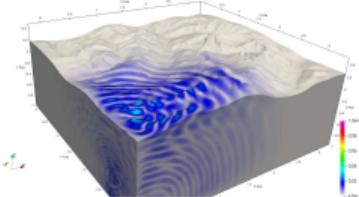
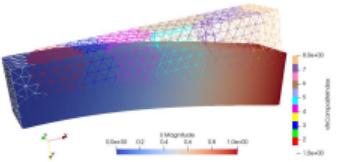
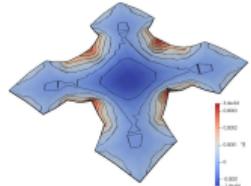


# ArcaneFEM: modules



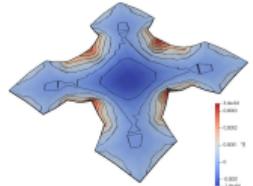


# ArcaneFEM: modules

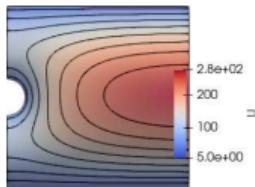




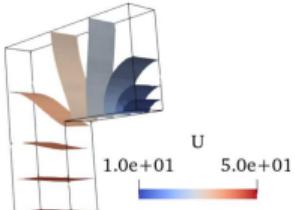
# ArcaneFEM: modules



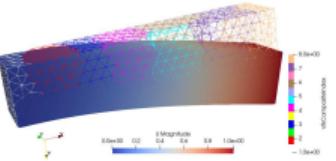
Bilaplacian  
Steady-Vector(2)-3D



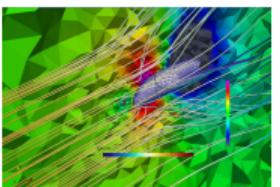
Fourier (GPU)  
Steady-Scalar-3D



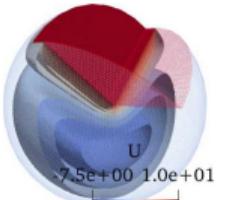
Laplace (GPU)  
Steady-Scalar-3D



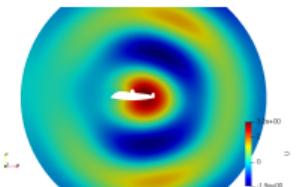
Elasticity (GPU)  
Steady-Vector(3)-3D



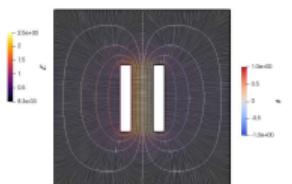
Aerodynamics  
Steady-Scalar-3D



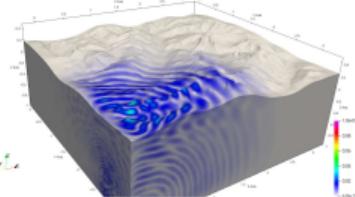
Poisson (GPU)  
Steady-Scalar-3D



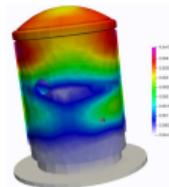
Acoustics  
Steady-Scalar-3D



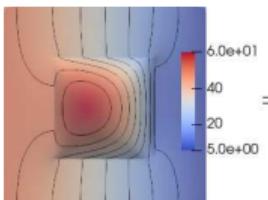
Electrostatic (GPU)  
Steady-Scalar-3D



Solid dynamics  
Transient-Vector(3)-4D



Elastodynamics  
Transient-Vector(3)-4D



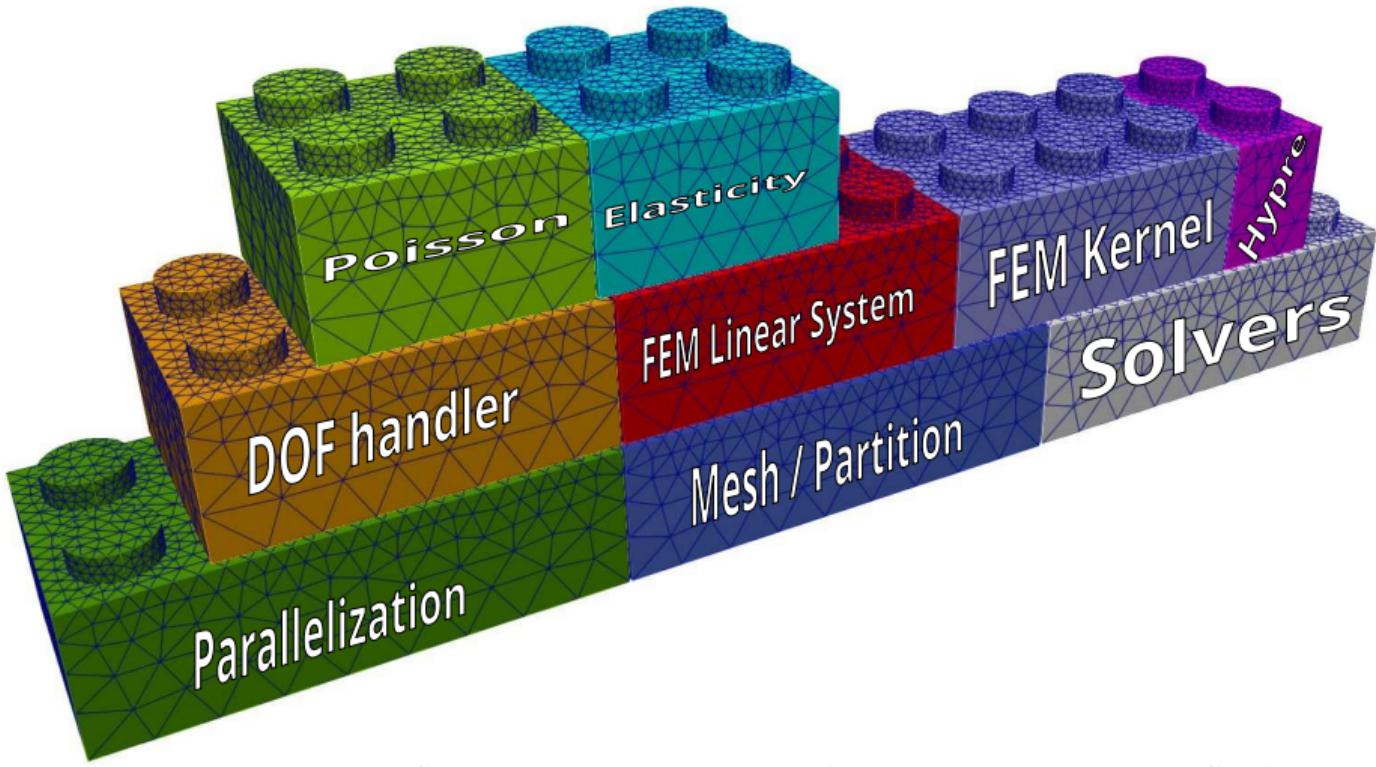
Heat transfer  
Transient-Scalar-4D



# ArcaneFEM: architecture

ArcaneFEM

Arcane

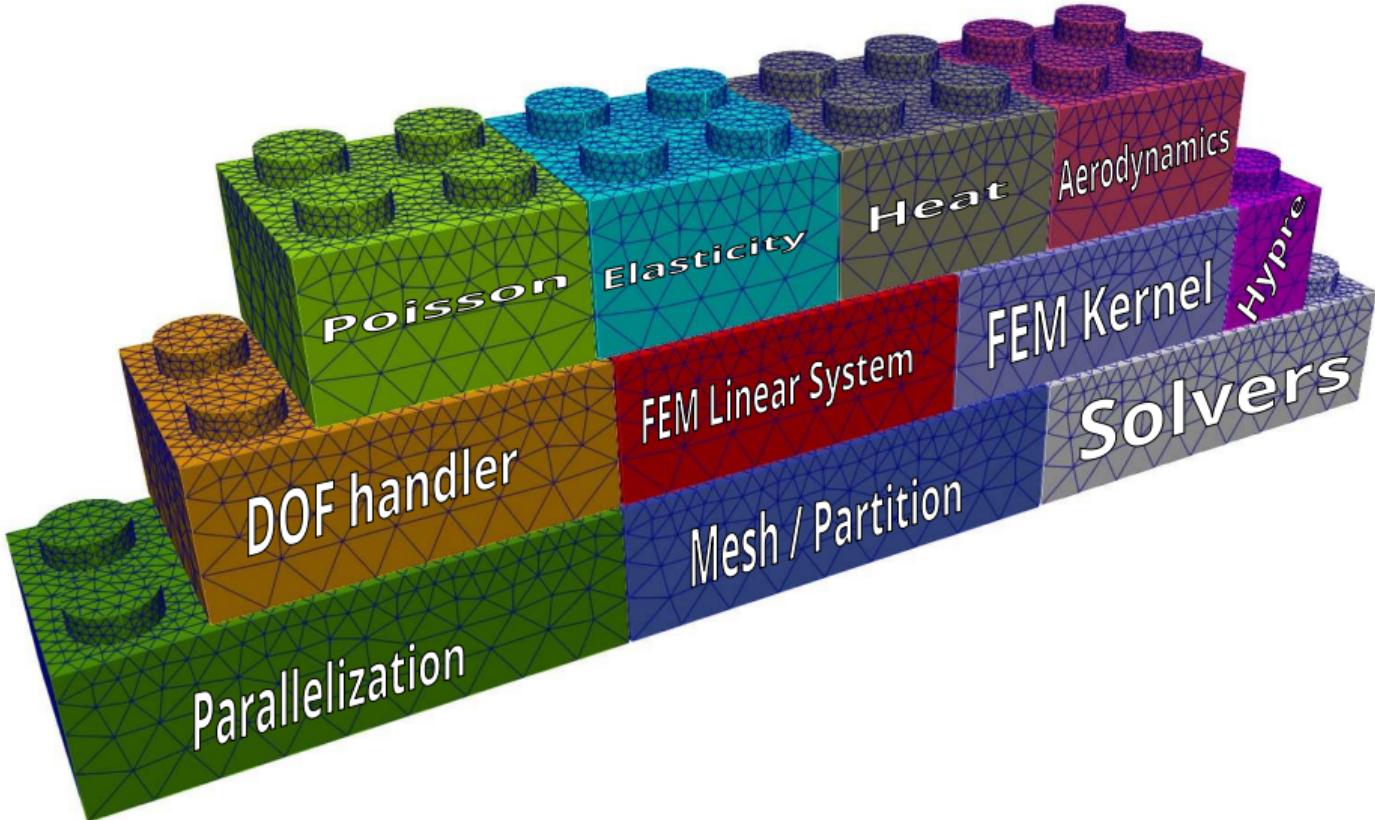




# ArcaneFEM: architecture

ArcaneFEM

Arcane



# Outline

## Introduction

## FEM on GPU

- Focus linear-system assembly
- Focus solving

## Results

- Benchmarks
- Large-scale test

## Key Takeaways & What's Next



## Outline



## Introduction

FEM on GPU

## Results

## Key Takeaways & What's Next

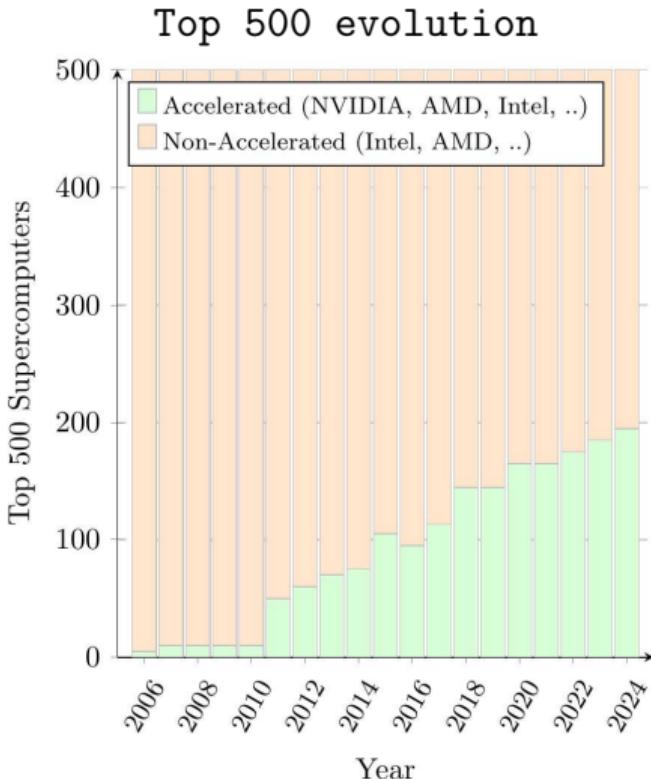


Motivation: why to evolve and use GPU's ?

- Earthquake simulations require solving billions of unknowns  $\mathcal{O}(10^9)$
  - Real time for **DA** / **ML**
  - Road to **Exascale** & green simulations
    - top 500: **193/500** are accelerated;  
top 10: **9/10** are accelerated
    - **70%** of top HPC applications are accelerated
    - 2026 French SC **Alice Recoque**

*“It is not the strongest of the species codes that survive, not the most intelligent that survives. It is the one that is most acceptable to change”*

Charles Darwin



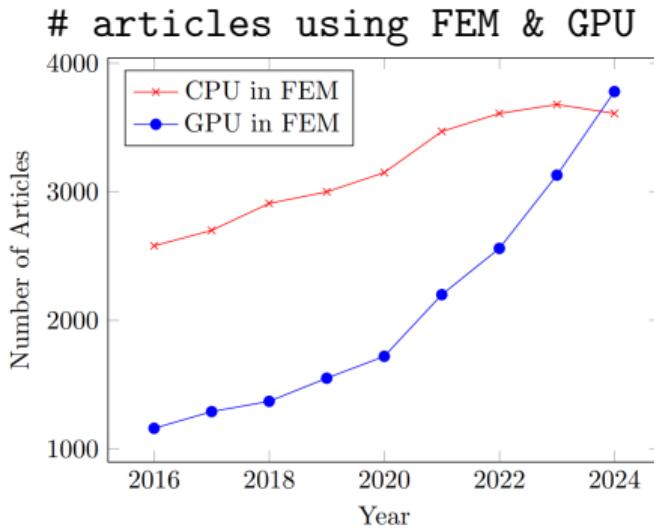


Motivation: why to evolve and use GPU's ?

- Earthquake simulations require solving billions of unknowns  $\mathcal{O}(10^9)$
  - Real time for **DA** / **ML**
  - Road to **Exascale** & green simulations
    - top 500: **193/500** are accelerated;  
top 10: **9/10** are accelerated
    - **70%** of top HPC applications are accelerated
    - 2026 French SC **Alice Recoque**
  - FEM is embracing GPU

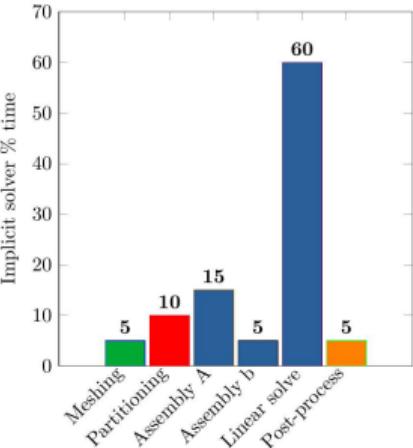
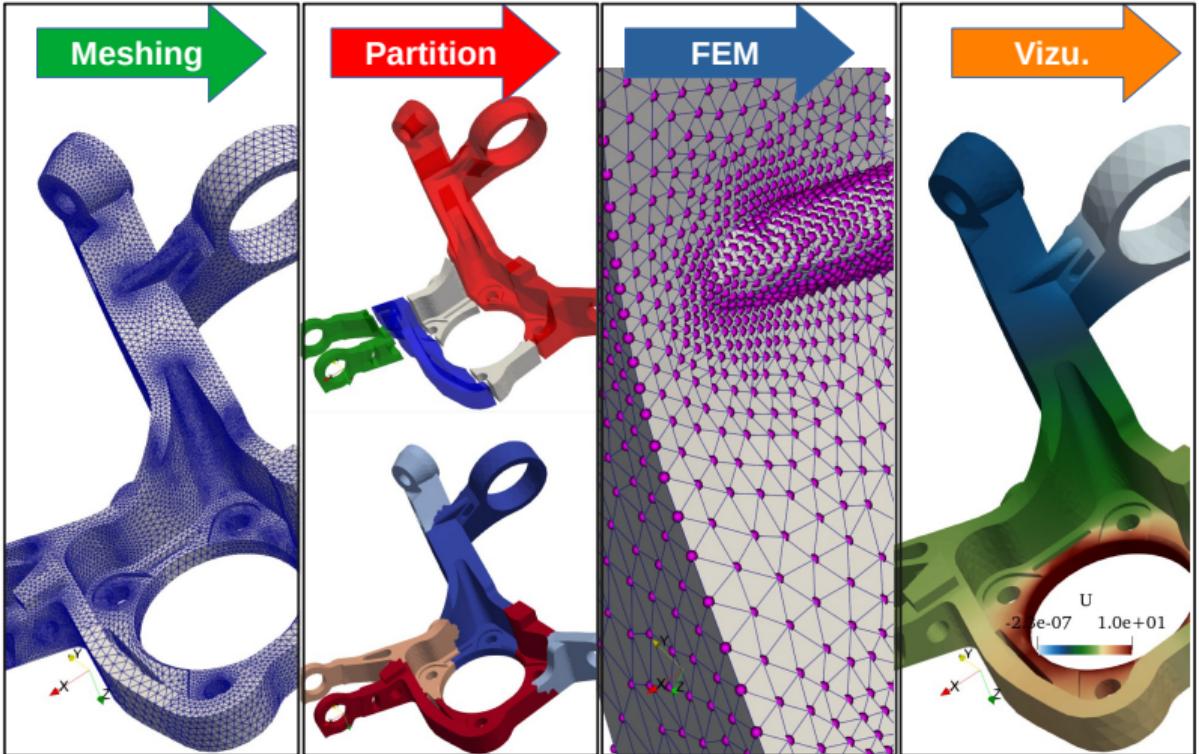
*“It is not the strongest of the species codes that survive, not the most intelligent that survives. It is the one that is most acceptable to change”*

Charles Darwin



## GPGPU for full FEM

A CPU-GPU simulation process in ArcaneFEM



- Full FEM on GPU
    - Assembly  $\mathbf{A}$ ,  $\mathbf{b}$
    - Solve
    - $$\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$$
  - Only Solving issues:
    - Data movement CPU-GPU
    - Low # MPI PIU

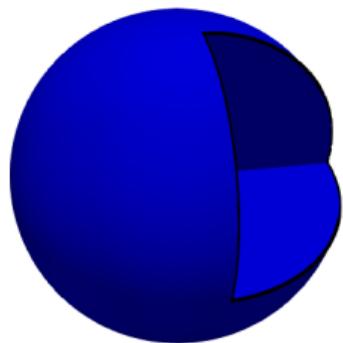


# Strategy for GPGPU-CPU simulations

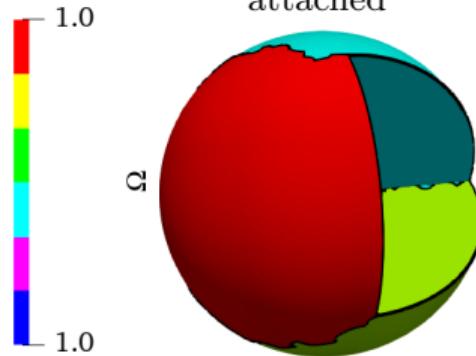
- Domain Decomposition paradigm for multi-GPU

- Attach MPI to GPU

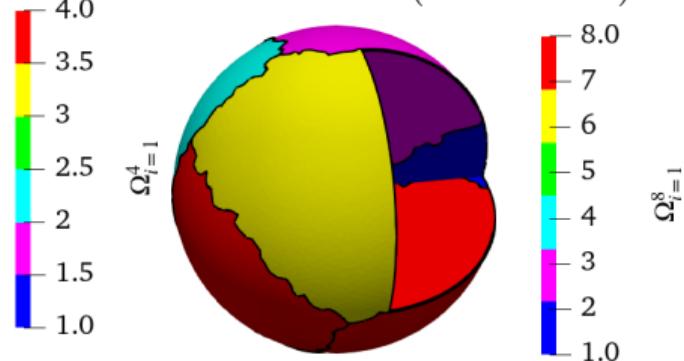
- mono 1 GPU simulation
- No partitioning, no MPI



- multi 4 GPU simulation
- 4 partitions, 4 MPI procs. attached



- multi 4 GPU-CPU simulation
- 8 partitions, 8 MPI procs. attached (batches of 2)



# Outline

Introduction

FEM on GPU

Results

Key Takeaways & What's Next



Introduction  
oooo

FEM on GPU  
●oooooooo

Results  
oooooooooooo

Conclusion  
oo

8/26



# FEM assembly: classic Galerkin

- Works for CPU, sequential / parallel
- Pitfall : GPGPU race-conditions / race-hazard prone

$$\int_{\Omega^h} \nabla u \cdot \nabla v = \int_{\partial\Omega_N^h} \nabla u \cdot \mathbf{n} + \int_{\Omega^h} fv \quad \& \quad Dirichlet$$

## Implementation classic

$$\int_{\Omega^h} \nabla u \cdot \nabla v \mapsto \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \mapsto \mathbf{A}$$

$$\int_{\partial\Omega_N^h} \nabla u \cdot \mathbf{n} + \int_{\Omega^h} fv \mapsto \int_{\mathcal{E}_i} \nabla u \cdot \mathbf{n} + \int_{\mathcal{T}_i} fv \mapsto \mathbf{b}$$





# FEM assembly: race-hazard demonstration

Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



Mesh

|   | 0 | 1 | 2 | 3 | 4 |
|---|---|---|---|---|---|
| 0 | ■ |   |   |   |   |
| 1 |   | ■ |   |   |   |
| 2 |   |   | ■ |   |   |
| 3 |   |   |   | ■ |   |
| 4 |   |   |   |   | ■ |

Matrix



# FEM assembly: race-hazard demonstration

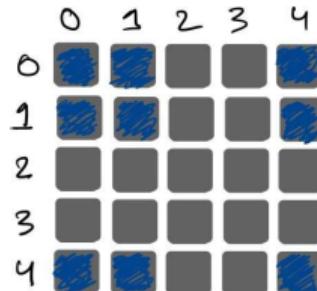
Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



Element Stiffness  
Matrix



Matrix



# FEM assembly: race-hazard demonstration

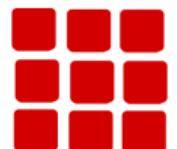
Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \mapsto \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \mapsto \mathbf{A}$$

$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



Mesh



Element Stiffness Matrix

|   |      |     |     |      |
|---|------|-----|-----|------|
| 0 | 1    | 2   | 3   | 4    |
| 0 | blue |     |     |      |
| 1 | red  | red |     |      |
| 2 |      | red | red |      |
| 3 |      |     | red |      |
| 4 | blue | red | red | blue |

Matrix





# FEM assembly: race-hazard demonstration

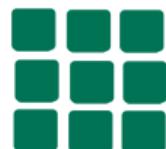
Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

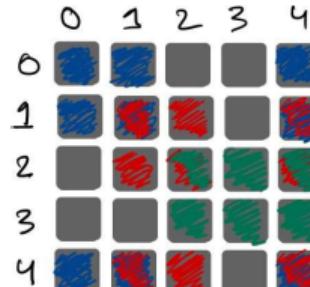
$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



Mesh



Element Stiffness Matrix



Matrix





# FEM assembly: race-hazard demonstration

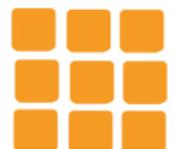
Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

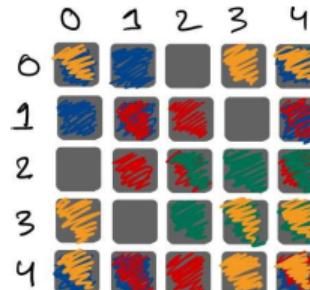
$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



Mesh



Element Stiffness Matrix



Matrix



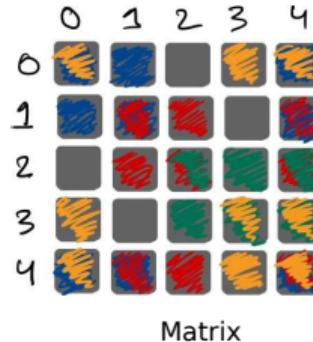


# FEM assembly: race-hazard demonstration

Global matrix via local stiffness matrix

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

$$A_{ij}^{(K)} = \sum_{p=1}^3 w_p \nabla \phi_i(x_p, y_p) \cdot \nabla \phi_j(x_p, y_p) |K|$$



```
ax::doAtomic<ax::eAtomicOperation::Add>(in_out_val_csr(begin), v);
```





# FEM assembly: Atomic Free Node-Wise Assembly

## Implementation classic

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{T}_i} \nabla u \cdot \nabla v \quad \mapsto \quad \mathbf{A}$$

$$\int_{\partial\Omega_N^h} \nabla u \cdot \mathbf{n} + \int_{\Omega^h} fv \quad \mapsto \quad \int_{\mathcal{E}_i} \nabla u \cdot \mathbf{n} + \int_{\mathcal{T}_i} fv \quad \mapsto \quad \mathbf{b}$$

## Node-Wise implementation

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \quad \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\int_{\partial\Omega_N^h} \nabla u \cdot \mathbf{n} + \int_{\Omega^h} fv \quad \mapsto \quad \int_{\mathcal{P}_i \in \partial\Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \partial\Omega^h} (\cdot) + \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{b}$$



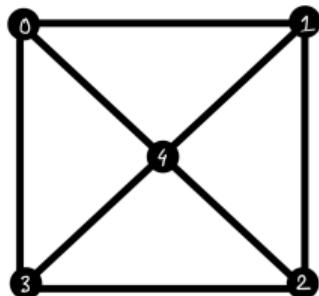


# Atomic Free Node-Wise Assembly: demonstration

Global matrix via **local stiffness vectors**

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\{A_i\} = \sum_{K \in \mathcal{K}_I} \sum_{p=1}^3 w_p (\nabla \phi_i(x_I, y_I) \otimes \nabla \phi_j(x_p, y_p)) |K|$$



Mesh

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 |
| 0 |   |   |   |   |
| 1 |   |   |   |   |
| 2 |   |   |   |   |
| 3 |   |   |   |   |
| 4 |   |   |   |   |

Matrix

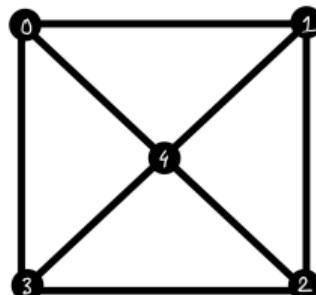


# Atomic Free Node-Wise Assembly: demonstration

Global matrix via **local stiffness vectors**

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\{A_i\} = \sum_{K \in \mathcal{K}_I} \sum_{p=1}^3 w_p (\nabla \phi_i(x_I, y_I) \otimes \nabla \phi_j(x_p, y_p)) |K|$$



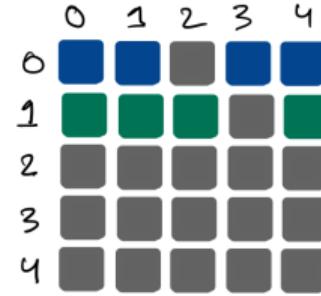
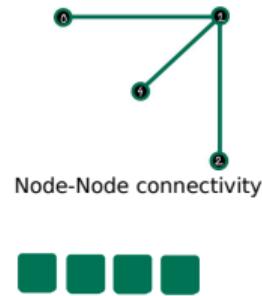
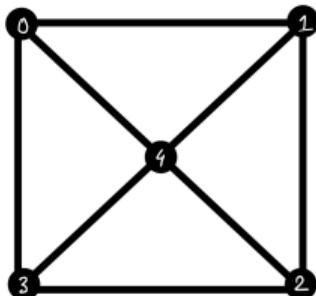


# Atomic Free Node-Wise Assembly: demonstration

Global matrix via **local stiffness vectors**

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\{A_i\} = \sum_{K \in \mathcal{K}_I} \sum_{p=1}^3 w_p (\nabla \phi_i(x_I, y_I) \otimes \nabla \phi_j(x_p, y_p)) |K|$$



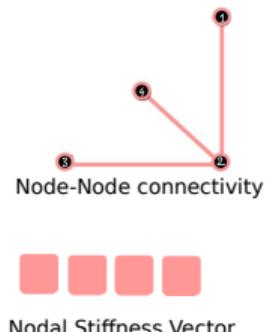
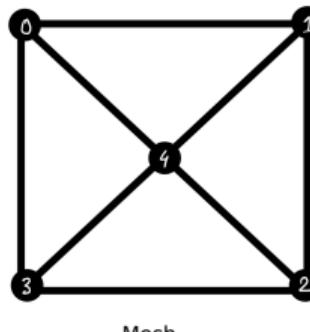


# Atomic Free Node-Wise Assembly: demonstration

Global matrix via **local stiffness vectors**

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\{A_i\} = \sum_{K \in \mathcal{K}_I} \sum_{p=1}^3 w_p (\nabla \phi_i(x_I, y_I) \otimes \nabla \phi_j(x_p, y_p)) |K|$$



| 0 | 1          | 2          | 3         | 4          |
|---|------------|------------|-----------|------------|
| 0 | Dark Blue  | Dark Blue  | Grey      | Dark Blue  |
| 1 | Dark Green | Dark Green | Grey      | Dark Green |
| 2 | Grey       | Light Red  | Light Red | Light Red  |
| 3 | Grey       | Light Red  | Light Red | Light Red  |
| 4 | Grey       | Light Red  | Light Red | Light Red  |

Matrix



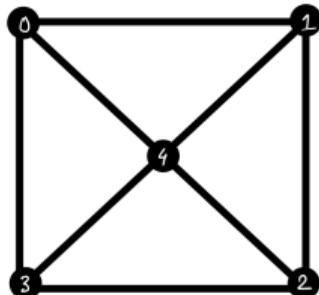


# Atomic Free Node-Wise Assembly: demonstration

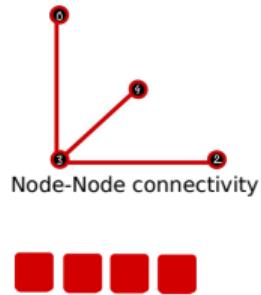
Global matrix via **local stiffness vectors**

$$\int_{\Omega^h} \nabla u \cdot \nabla v \quad \mapsto \int_{\mathcal{P}_i \in \Omega^h} (\cdot) + \int_{\mathcal{E}_i \in \Omega^h} (\cdot) \quad \mapsto \quad \mathbf{A}$$

$$\{A_i\} = \sum_{K \in \mathcal{K}_I} \sum_{p=1}^3 w_p (\nabla \phi_i(x_I, y_I) \otimes \nabla \phi_j(x_p, y_p)) |K|$$



Mesh



Nodal Stiffness Vector

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 |
| 0 | ■ | ■ | ■ | ■ |
| 1 | ■ | ■ | ■ | ■ |
| 2 | ■ | ■ | ■ | ■ |
| 3 | ■ | ■ | ■ | ■ |
| 4 | ■ | ■ | ■ | ■ |

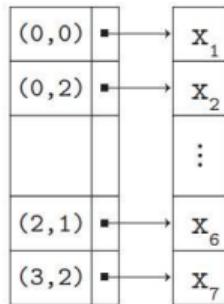
Matrix



# FEM matrices: various data-structures in ArcaneFEM

|       |       |       |       |
|-------|-------|-------|-------|
| $x_1$ | 0     | $x_2$ | $x_3$ |
| 0     | 0     | $x_4$ | 0     |
| $x_5$ | $x_6$ | 0     | 0     |
| 0     | 0     | $x_7$ | 0     |

## Dictionary Of Keys (DOK)



C++ `std::unordered_map`

- CPU only, slow search
- random memory access
- Easy to implement, scales well

## (Block) Compress Sparse Row ((B)CSR)

|      |       |       |       |       |       |       |       |
|------|-------|-------|-------|-------|-------|-------|-------|
| row: | 0     | 3     | 4     | 6     |       |       |       |
| col: | 0     | 2     | 3     | 2     | 0     | 1     | 2     |
| val: | $x_1$ | $x_2$ | $x_3$ | $x_4$ | $x_5$ | $x_6$ | $x_7$ |

matrix with 3 arrays + `row` is compressed

- GPU & CPU, fast search

## COOrdinate list (COO)

|      |       |       |       |       |       |       |       |
|------|-------|-------|-------|-------|-------|-------|-------|
| row: | 0     | 0     | 0     | 1     | 2     | 3     | 3     |
| col: | 0     | 2     | 3     | 2     | 0     | 1     | 2     |
| val: | $x_1$ | $x_2$ | $x_3$ | $x_4$ | $x_5$ | $x_6$ | $x_7$ |

matrix with 3 arrays

- GPU & CPU, fast search



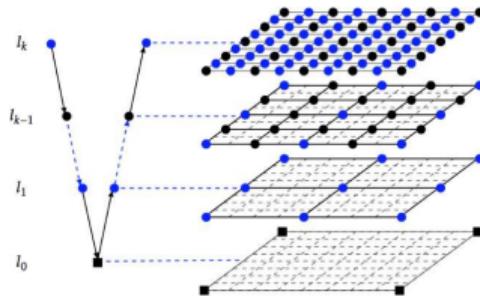
# FEM solving: multi-GPU solving via Hypre

## What:

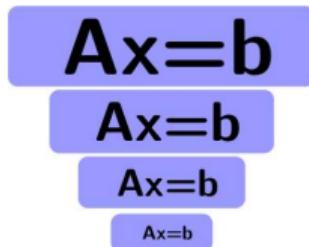
- Use hierarchy discretizations
  - meshes  $\mapsto$  GMG, Matrices  $\mapsto$  AMG
- Restrict and interpolate cycle
- Equations of interest elliptic : multigrid precod.
- AMG is used as precond. to Krylov solver CG

## How:

- Hypre open-source : BoomerAMG
  - + multi-GPU support (NVIDIA, AMD (2.28))
  - + Coarsner on GPU
  - + Smoothenr on GPU
  - + Interpolation Kernel on GPU
  - - CPU-GPU not handled
  - - COO format not handled



- Geometric multigrid



- Algebraic multigrid



# Outline

Introduction

FEM on GPU

Results

Key Takeaways & What's Next



Introduction  
oooo

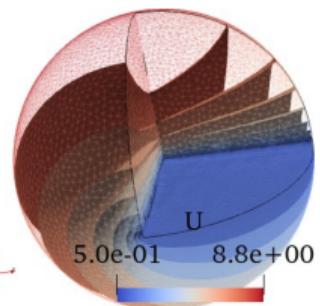
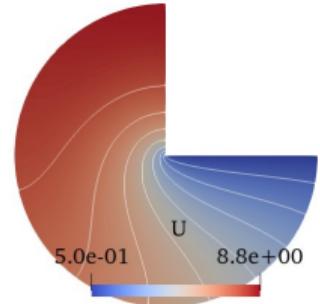
FEM on GPU  
oooooooo

Results  
●○○○○○○○○○○

Conclusion  
oo

15 / 26

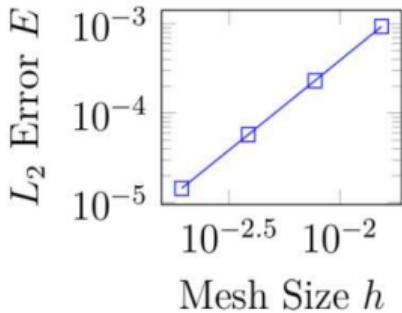
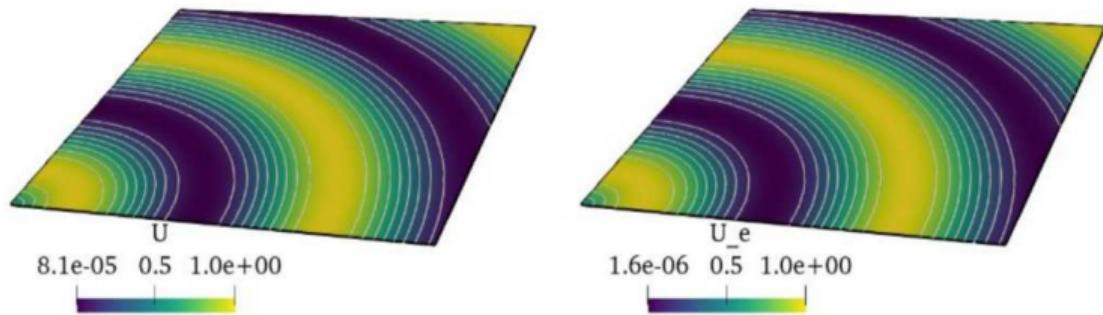
## Test Bench



## Poissons Problem 2D / 3D



# Solver verification via manufactured solution



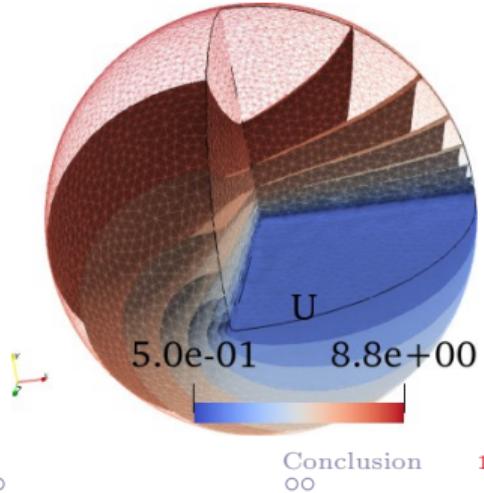
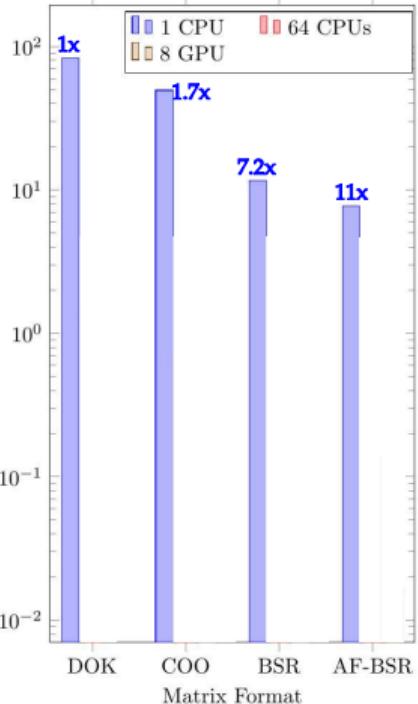
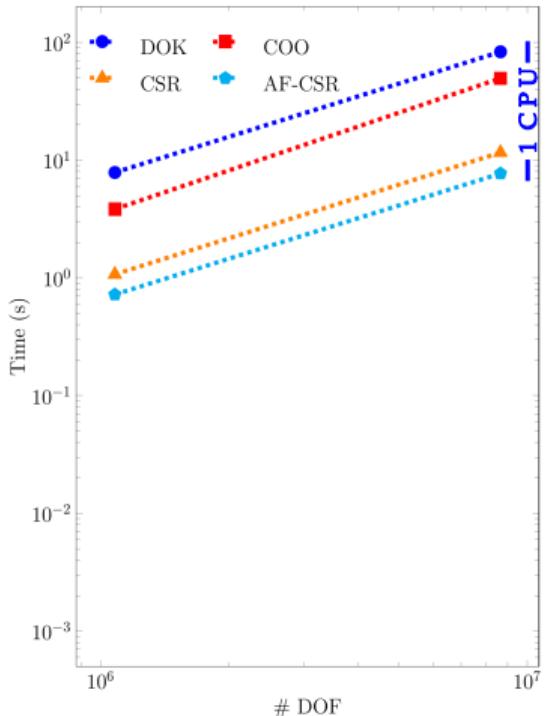
Comparison of FEM solution (left) from ArcaneFEM using the new GPU algorithm and the reference exact solution (right) from the following equation :

$$u_e(\boldsymbol{x}) = \frac{1}{2} + \frac{1}{2} \sin\left(\frac{\pi\sqrt{x^2 + y^2}}{\alpha}\right) \quad \forall \boldsymbol{x} \in \Omega$$



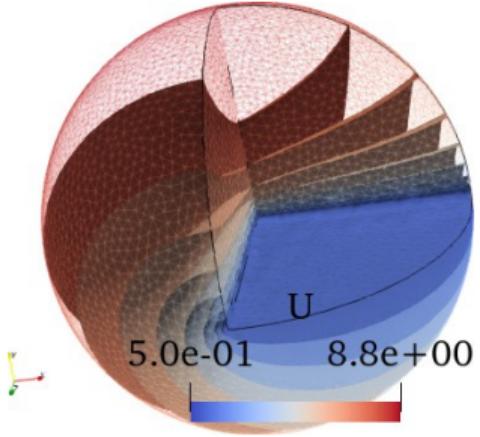
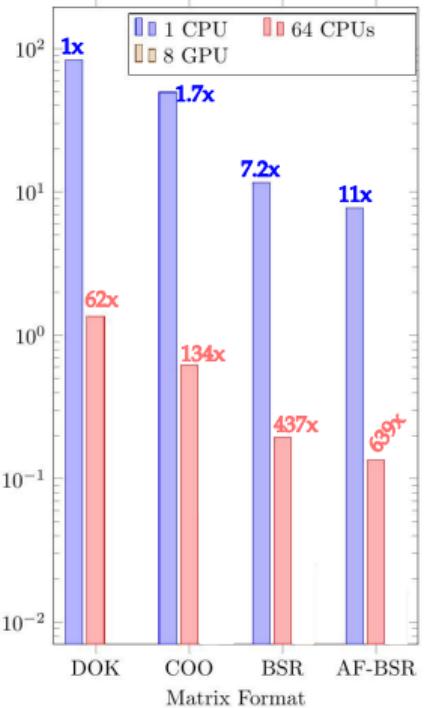
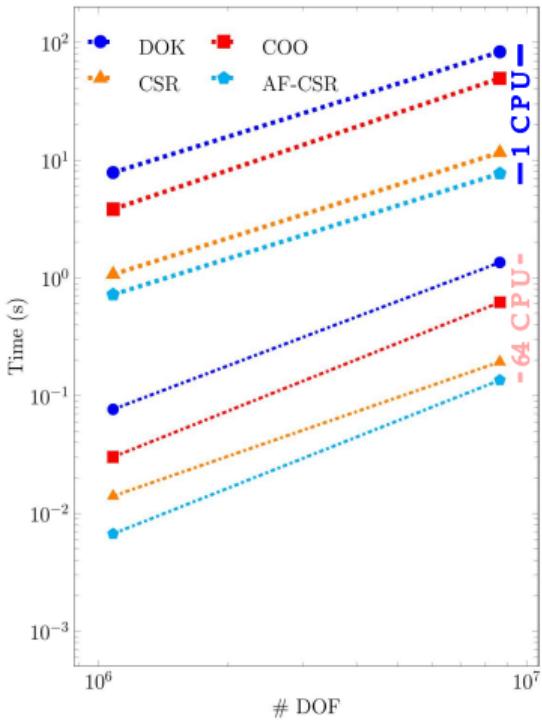


# Benchmark Supercomputer $\forall N_{\text{DOF}} = \{1\text{M}, 10\text{M}\}$



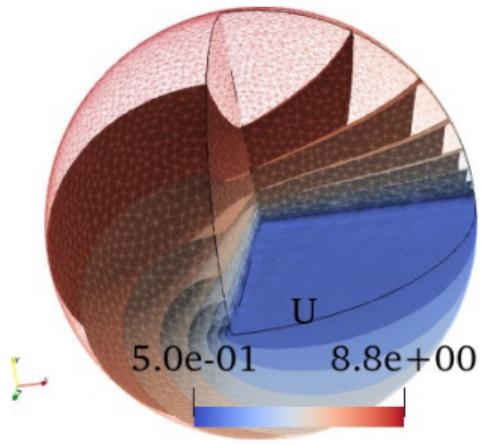
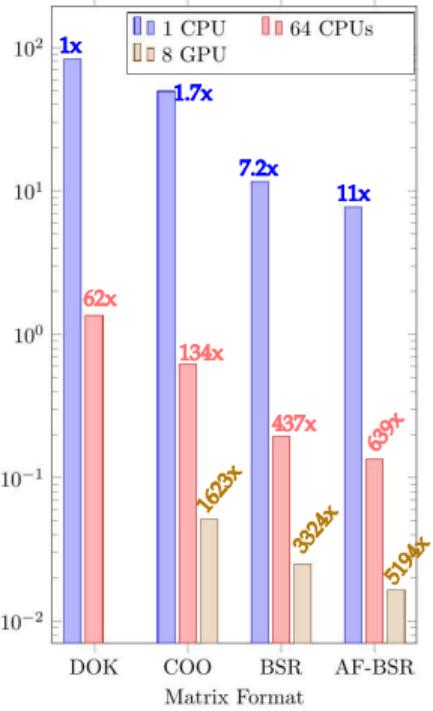
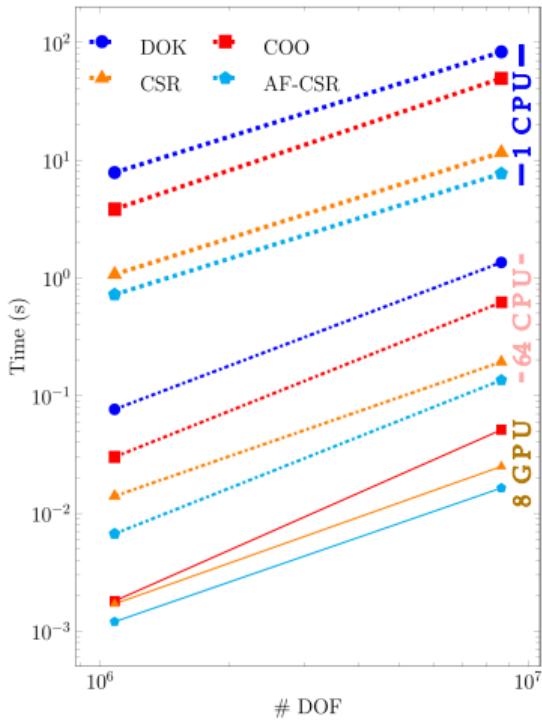


# Benchmark Supercomputer $\forall N_{\text{DOF}} = \{1\text{M}, 10\text{M}\}$



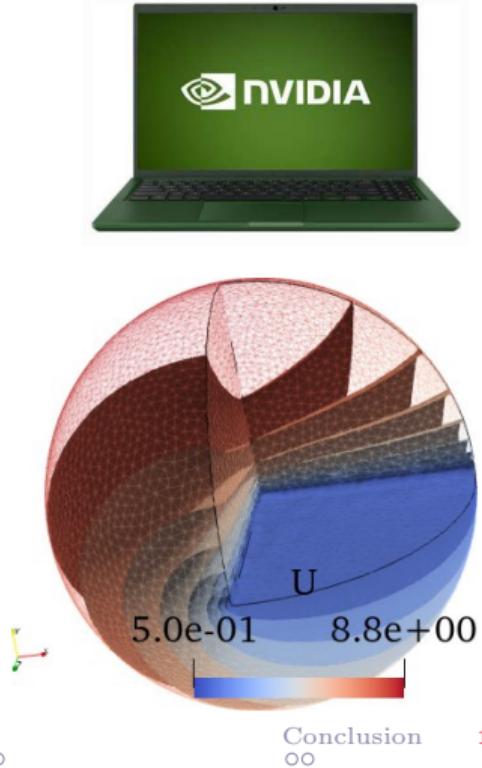
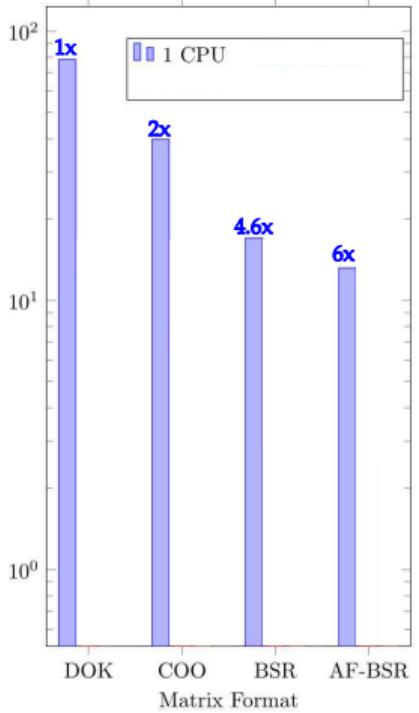
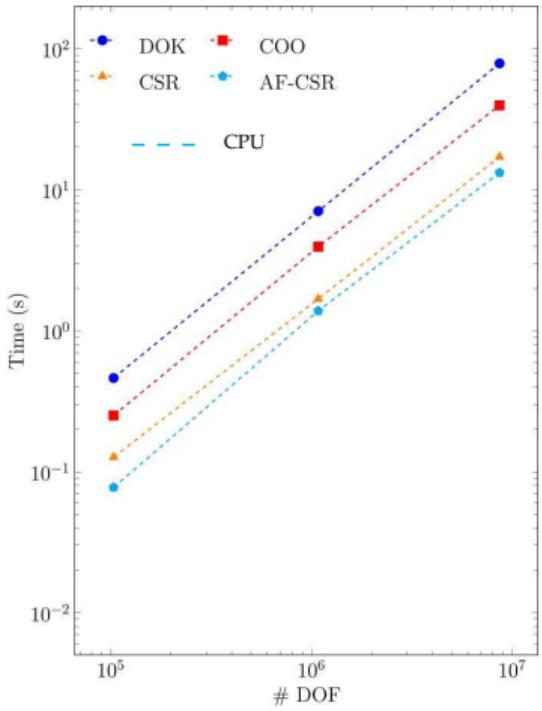


# Benchmark Supercomputer $\forall N_{\text{DOF}} = \{1\text{M}, 10\text{M}\}$



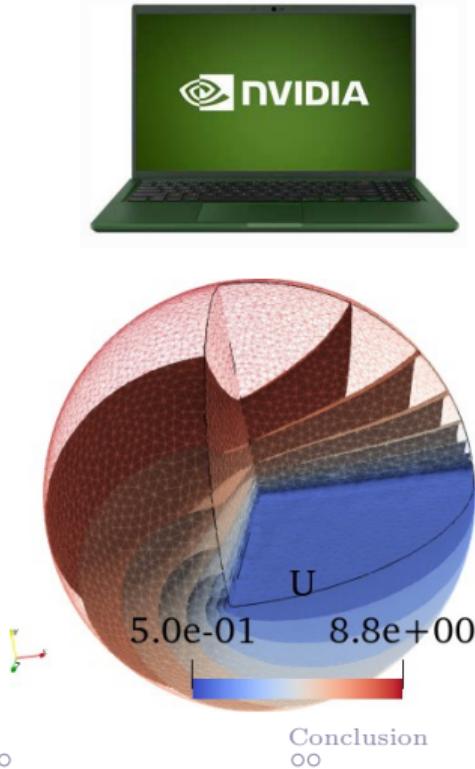
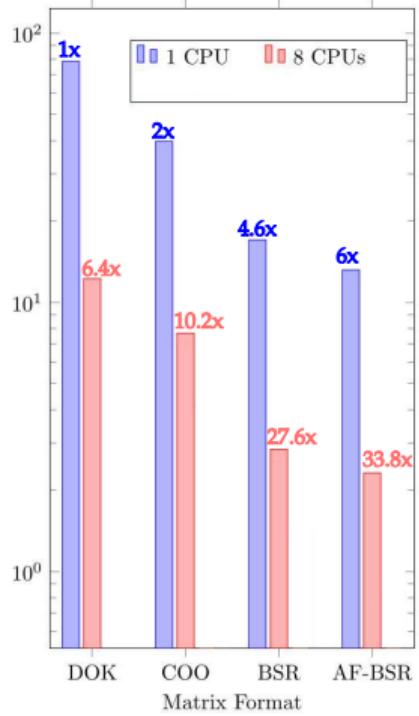
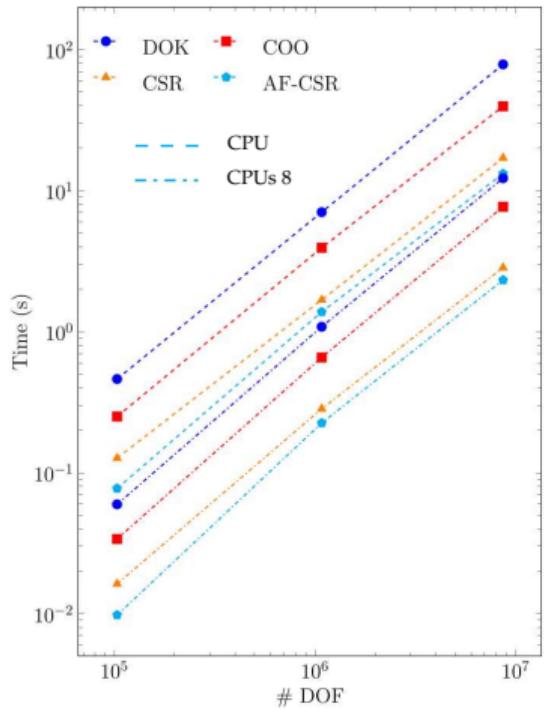


# Benchmark Laptop $\forall N_{\text{DOF}} = \{100\text{K}, 1\text{M}, 10\text{M}\}$



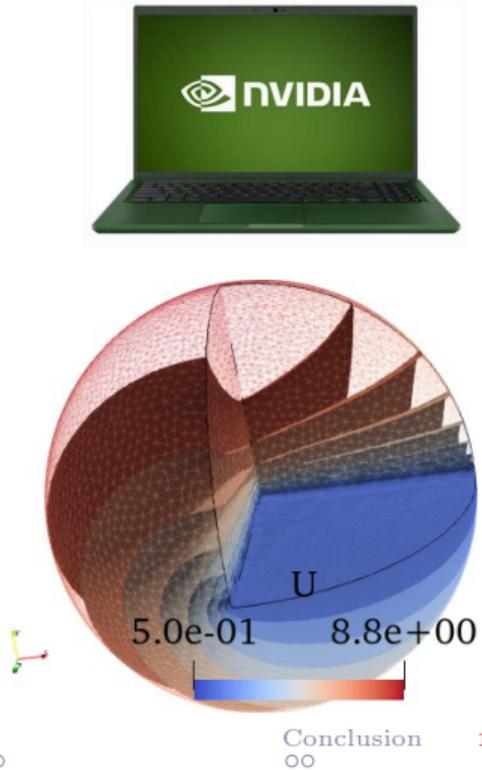
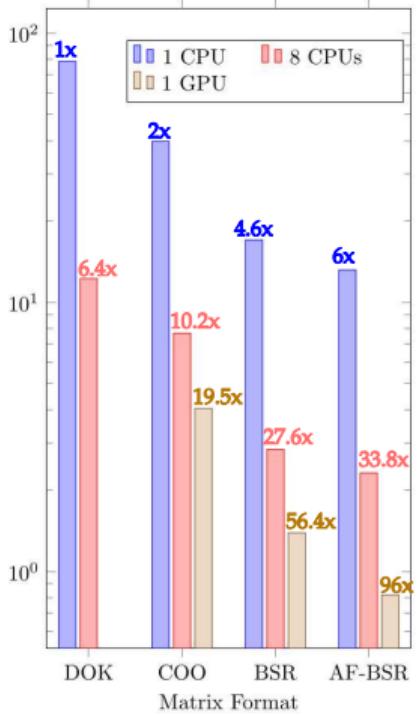
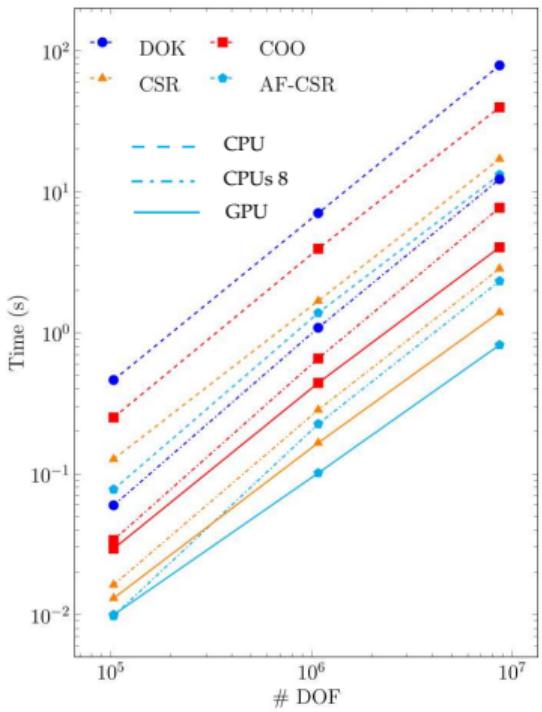


# Benchmark Laptop $\forall N_{\text{DOF}} = \{100\text{K}, 1\text{M}, 10\text{M}\}$



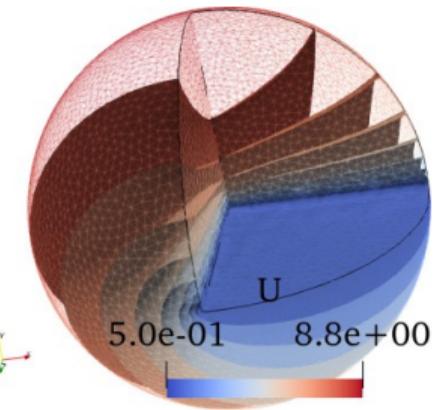
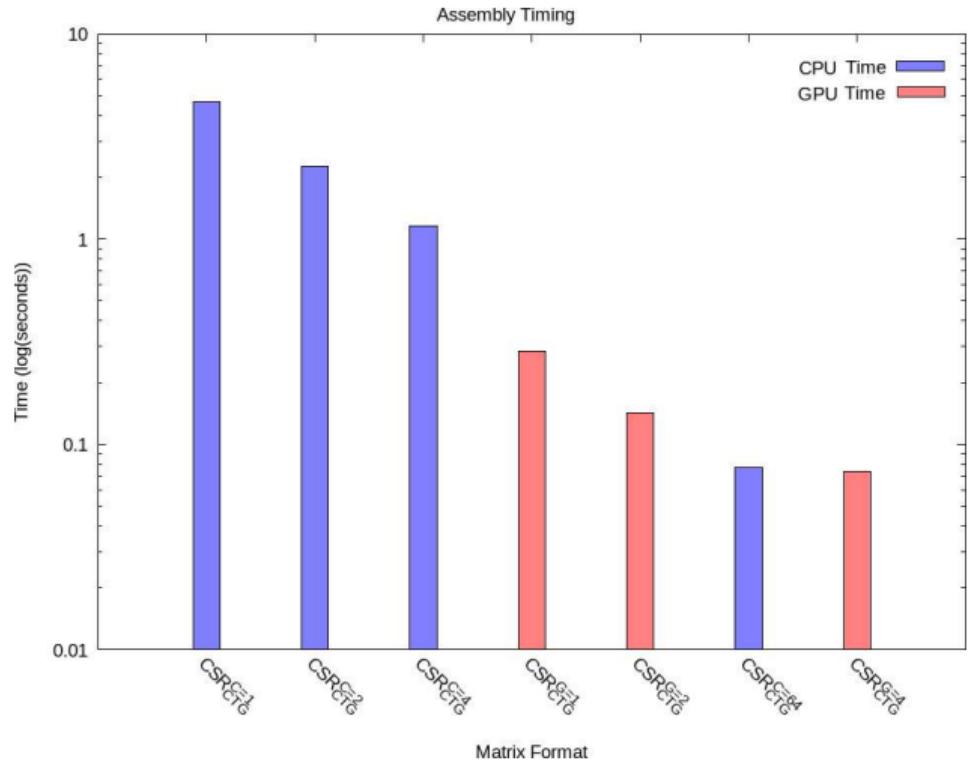


# Benchmark Laptop $\forall N_{\text{DOF}} = \{100K, 1M, 10M\}$

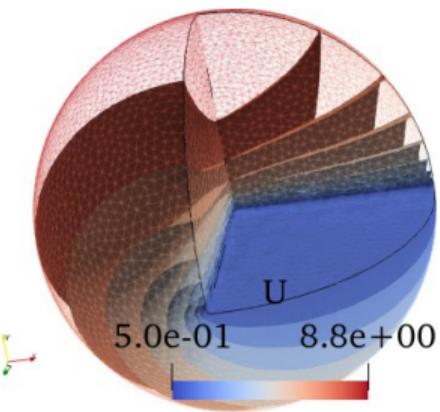
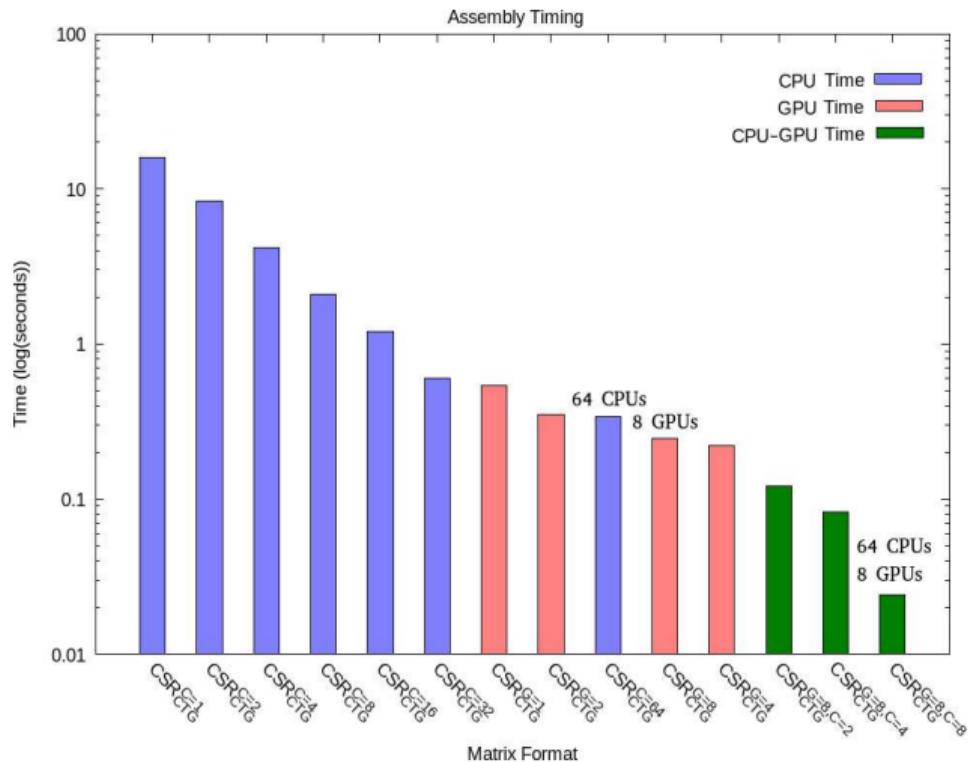




# Benchmark Supercomputers $N_{DOF} = 18M$

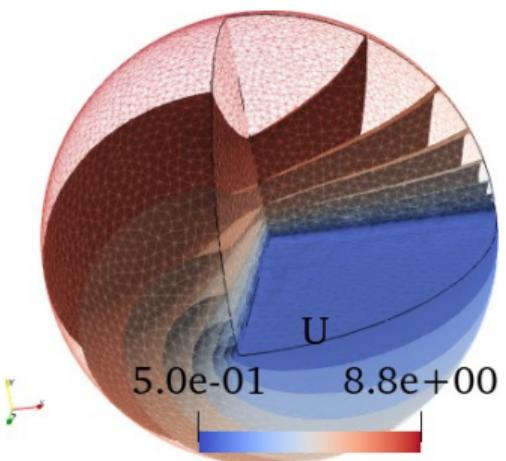
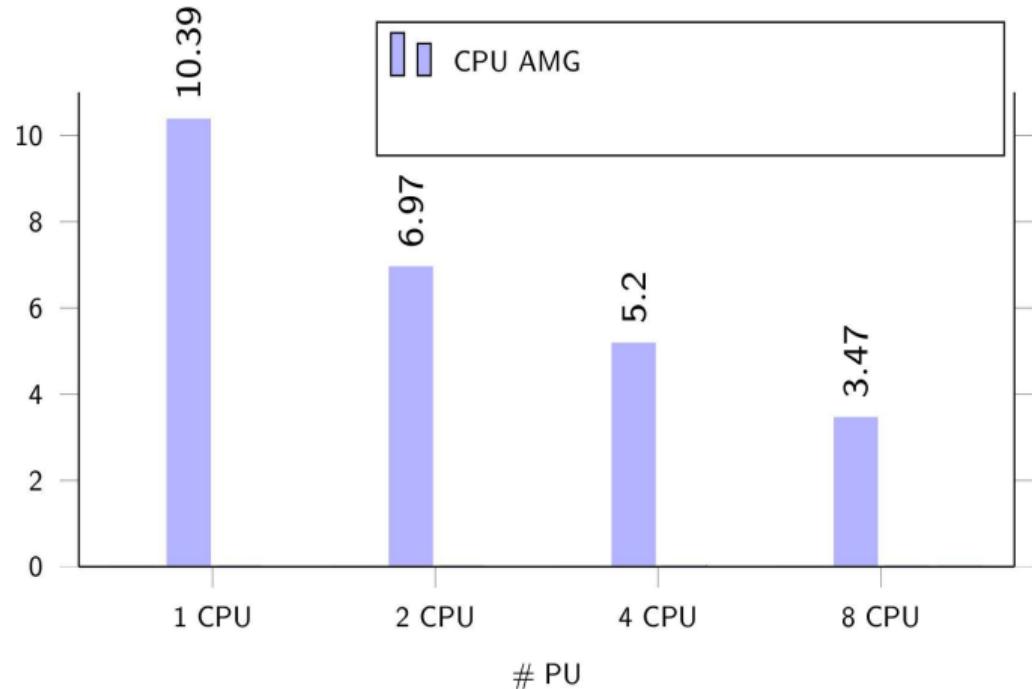


Benchmark Supercomputers  $N_{\text{DOF}} = 18\text{M}$



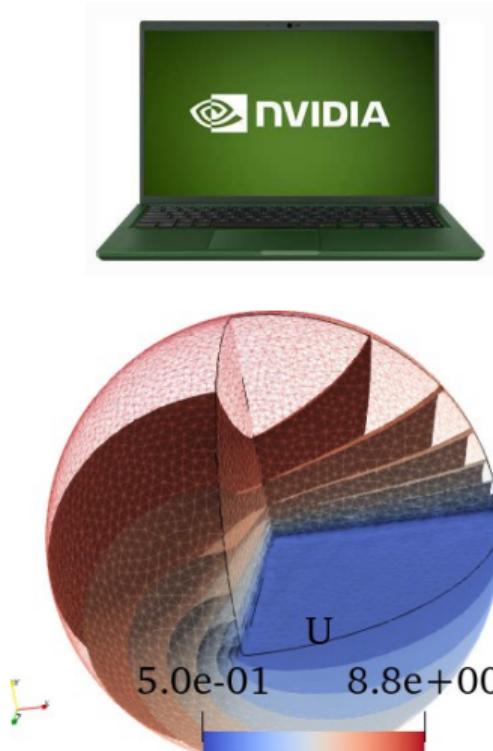
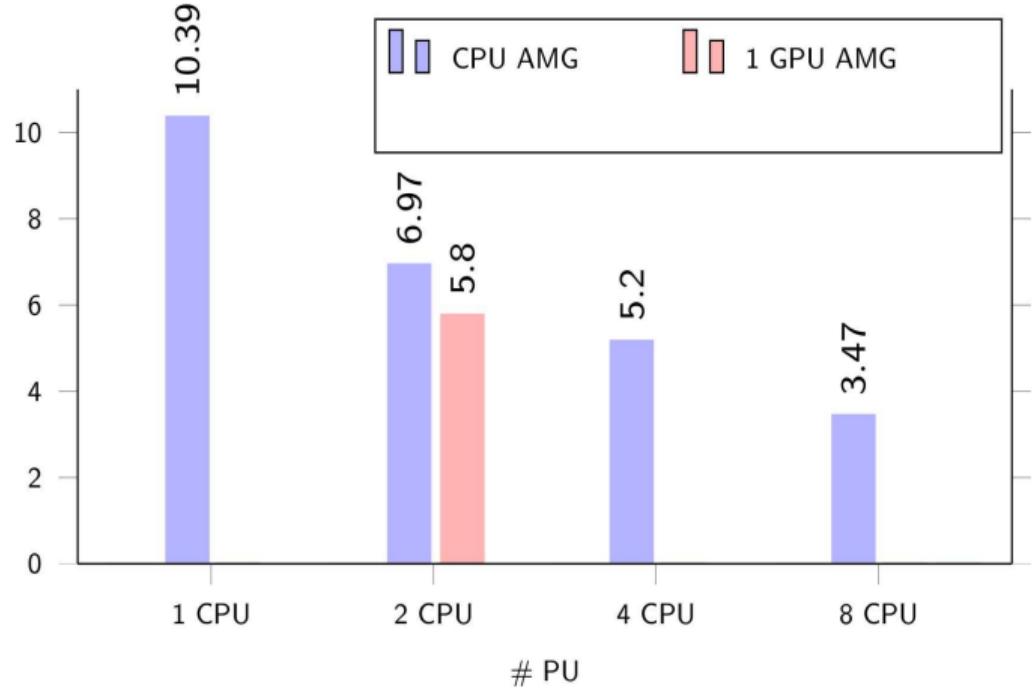


# Benchmark Laptop Preconditioner $N_{DOF} = 3.5M$



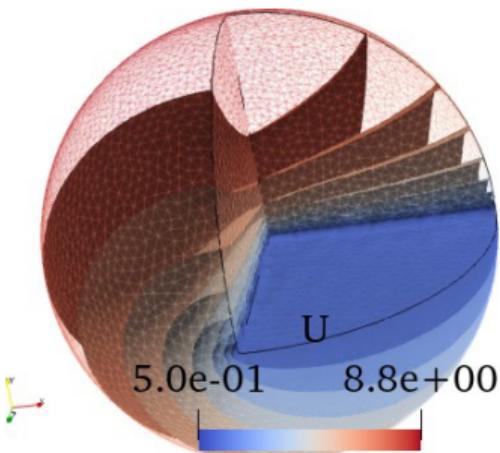
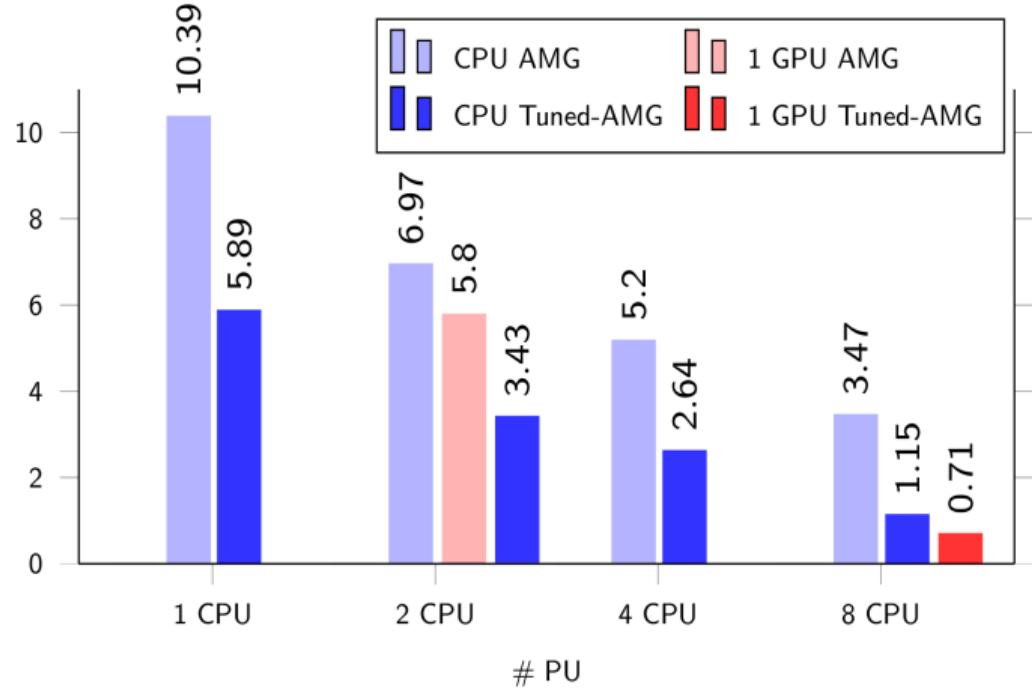


# Benchmark Laptop Preconditioner $N_{DOF} = 3.5M$



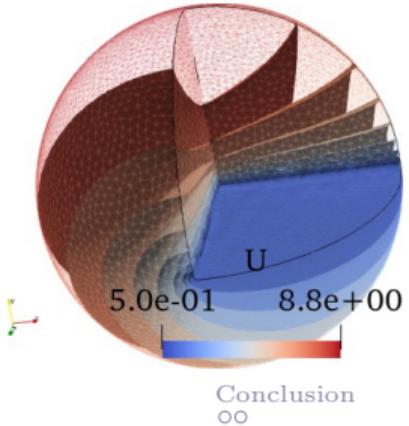
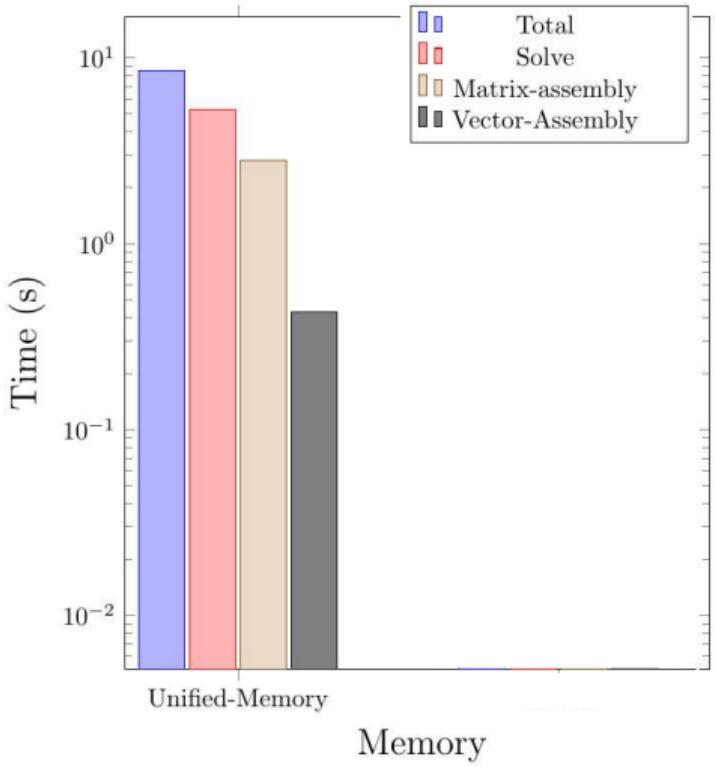


# Benchmark Laptop Preconditioner $N_{DOF} = 3.5M$



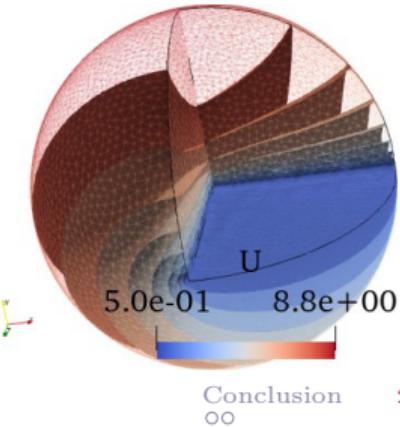
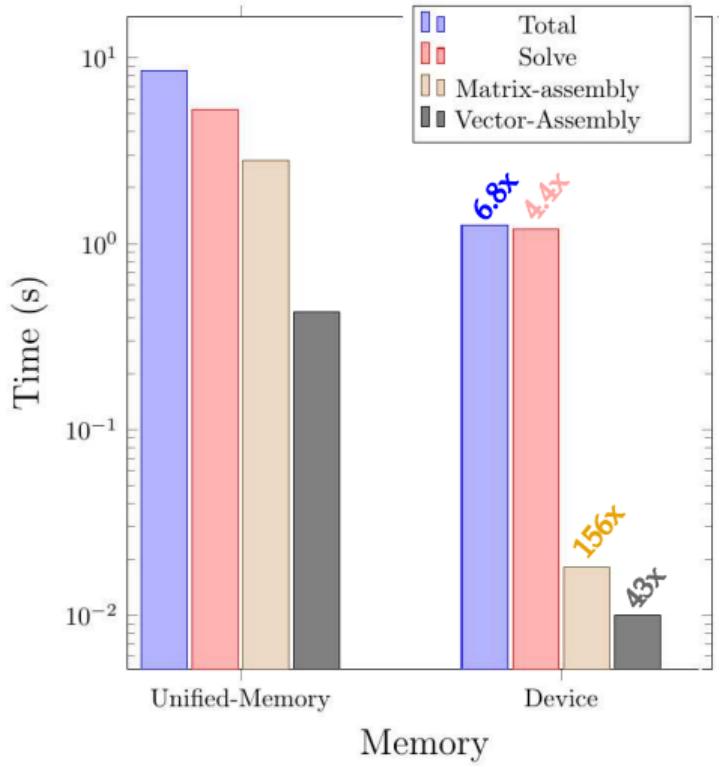


# Memory resource should we care ?



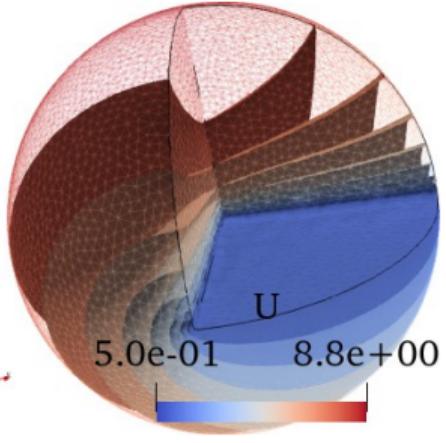
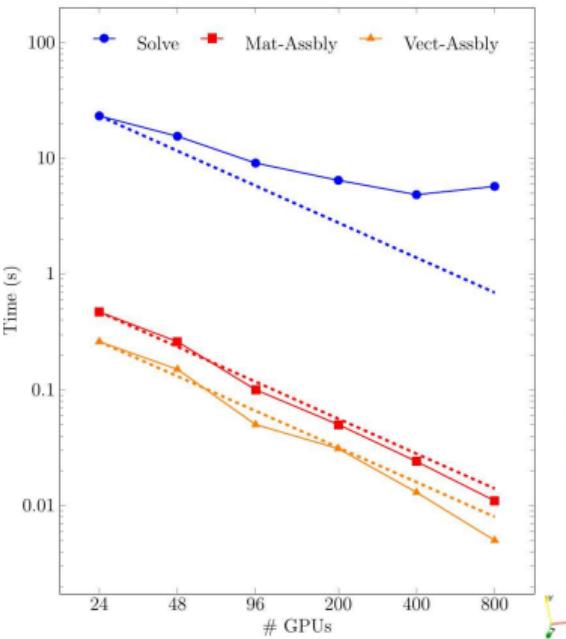
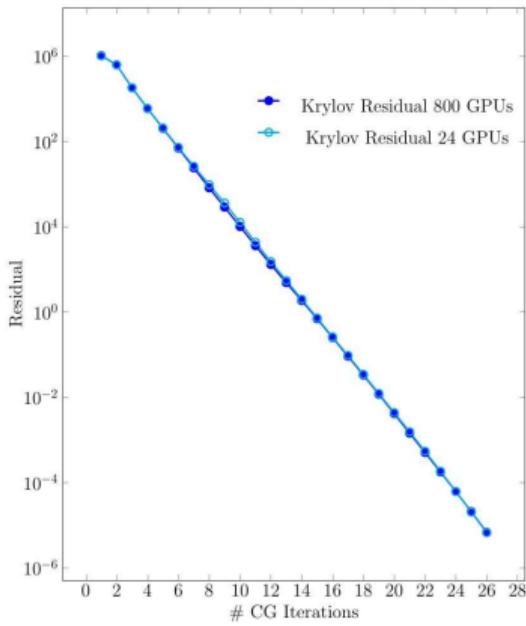


# Memory resource should we care ?



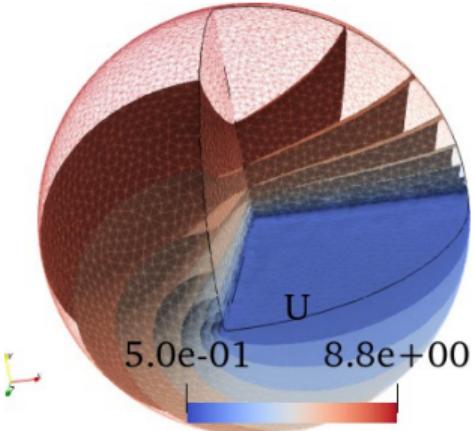
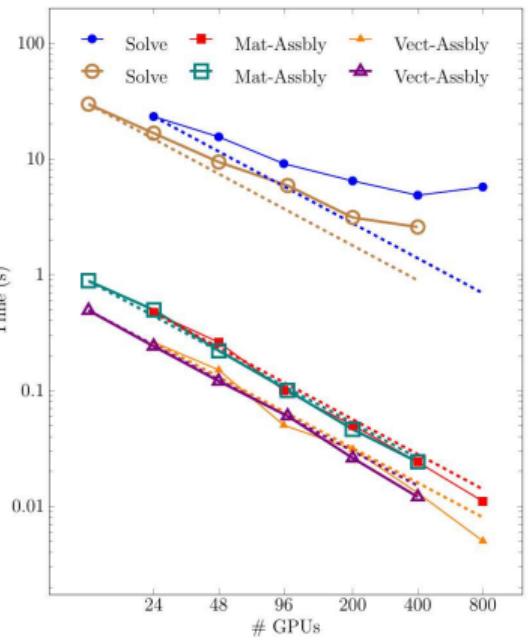
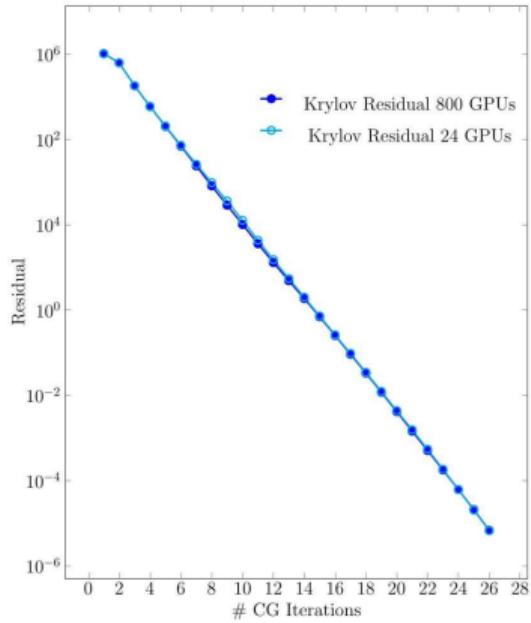


# Benchmark $N_{\text{DOF}} = 0.27\text{B}$ on 800 GPUs





# Benchmark $N_{DOF} = 0.27B$ on 800 GPUs





# 1B cell problem on 100 GPUs

## Mesh:

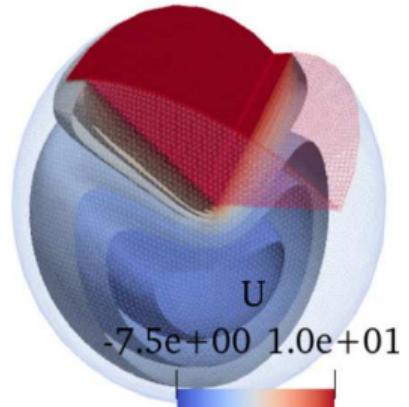
- 1 Billion Tetrahedrons
- KNL fat-node (128 Threads Gmsh)

## Partitioning:

- dual-level 100 MPI Mesh reading
- Parmetis

## Timing:

- Matrix assembly 0.0001 (s)
- Solving 210 (s) (**not tuned**)



# Outline

Introduction

FEM on GPU

Results

Key Takeaways & What's Next





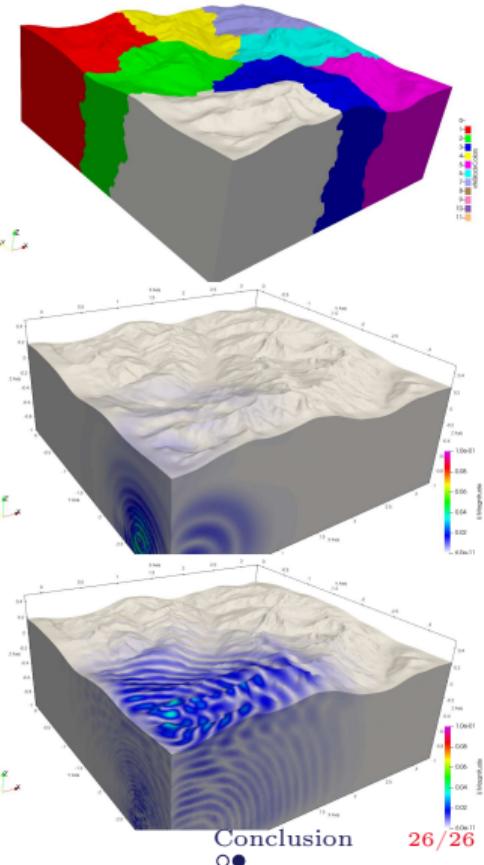
# Key Takeaways & What's Next

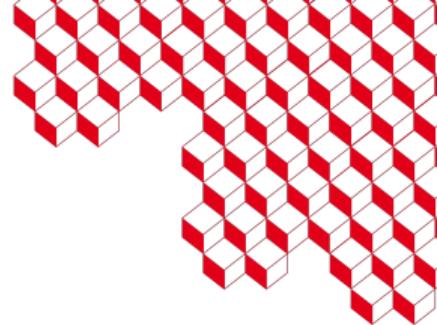
## Conclusion

- ArcaneFEM first works on **GPGPU & CPU-GPU** FEM solving
- **Matrix assembly** via porting approach : **Atomics**
- Atomic free assembly, using **node-wise** approach
- Benchmarks on GPUs: **MI250X, A100, & Laptop**

## Future Works

- GPU for all modules (Elastodynamic-Soildynamics)
- Focus on **non-linear** physics, **multi-physics**, **PASSMO** modules
- **top-ii-vol** integration: combined **mesher-partitioner** on GPUs
- APU's benchmarking: **MI300A, GH200**
- Linear system solve:
  - CPU = theoretical AMG ; GPU != theoretical AMG
  - AMG with **full CPU-GPU**
  - **Matrix free solvers**





ArcaneFEM: Scalable GPGPU Implicit FEM Solver for Unstructured Meshes



# Backup1

## Weak form of Poisson's problem

$$\int_{\Omega} \nabla u \cdot \nabla v \, d\Omega = \int_{\Omega} f v \, d\Omega, \quad \forall v \in V.$$

Using basis functions  $\{\phi\}_{i=1}^3$

$$u_h = \sum_i U_i \phi_i, \quad v_h = \phi_j.$$

Substituting into the weak form:

$$\sum_i U_i \int_{\Omega} \nabla \phi_i \cdot \nabla \phi_j \, d\Omega = \int_{\Omega} f \phi_j \, d\Omega.$$

This leads to:

$$\mathbf{A}_{ij} = \int_{\Omega} \nabla \phi_i \cdot \nabla \phi_j \, d\Omega.$$

## Local Stiffness Matrix for Element $K$

For each finite element  $K$ , the local stiffness matrix is:

$$A_{ij}^{(K)} = \int_K \nabla \phi_i \cdot \nabla \phi_j \, dK.$$

For TRIA3  $\{i\}_{k=1}^3 \{j\}_{k=1}^3$  element:

$$A^{(K)} = \begin{bmatrix} a_{11}^{(K)} & a_{12}^{(K)} & a_{13}^{(K)} \\ a_{21}^{(K)} & a_{22}^{(K)} & a_{23}^{(K)} \\ a_{31}^{(K)} & a_{32}^{(K)} & a_{33}^{(K)} \end{bmatrix}.$$



# Backup1

```
14 /*-----*/
15 /**
16 * @brief Computes the element matrix for a triangular element (P1 FE).
17 *
18 * This function calculates the integral of the expression:
19 *  $\int \int (\frac{\partial u}{\partial x} \frac{\partial v}{\partial x} + \frac{\partial u}{\partial y} \frac{\partial v}{\partial y}) d\Omega$ 
20 *
21 * Steps involved:
22 * 1. Calculate the area of the triangle.
23 * 2. Compute the gradients of the shape functions.
24 * 3. Return  $\int \int (\frac{\partial u}{\partial x} \frac{\partial v}{\partial x} + \frac{\partial u}{\partial y} \frac{\partial v}{\partial y}) d\Omega$ .
25 */
26 /*-----*/
27
28 RealMatrix<3, 3> FemModule::computeElementMatrixTria3(Cell cell)
29 {
30     Real area = ArcaneFemFunctions::MeshOperation::computeAreaTria3(cell, m_node_coord);
31
32     Real3 dxU = ArcaneFemFunctions::FeOperation2D::computeGradientXTria3(cell, m_node_coord);
33     Real3 dyU = ArcaneFemFunctions::FeOperation2D::computeGradientYTri3(cell, m_node_coord);
34
35     return area * (dxU ^ dxU) + area * (dyU ^ dyU);
36 }
37
38 /*-----*/
39 /*-----*/
40 /*-----*/
41
42 ARCCORE_HOST_DEVICE RealMatrix<3, 3>
43 _computeElementMatrixTria3Gpu(CellLocalId cell_lid, const IndexedCellNodeConnectivityView& cn_cv, const ax::VariableNodeReal3InView& in_node_coord)
44 {
45     Real area = Arcane::FemUtils::Gpu::MeshOperation::computeAreaTria3(cell_lid, cn_cv, in_node_coord);
46
47     Real3 dxU = FemUtils::Gpu::FeOperation2D::computeGradientXTria3(cell_lid, cn_cv, in_node_coord);
48     Real3 dyU = FemUtils::Gpu::FeOperation2D::computeGradientYTri3(cell_lid, cn_cv, in_node_coord);
49
50     return area * (dxU ^ dxU) + area * (dyU ^ dyU);
51 }
52
53 /*-----*/
54 /*-----*/
55
56 ARCCORE_HOST_DEVICE RealMatrix<1, 3>
57 _computeElementVectorTria3Gpu(CellLocalId cell_lid, const IndexedCellNodeConnectivityView& cn_cv, const ax::VariableNodeReal3InView& in_node_coord, Int32 node_lid)
58 {
59     Real area = Arcane::FemUtils::Gpu::MeshOperation::computeAreaTria3(cell_lid, cn_cv, in_node_coord);
60
61     Real3 dxU = FemUtils::Gpu::FeOperation2D::computeGradientXTria3(cell_lid, cn_cv, in_node_coord);
62     Real3 dyU = FemUtils::Gpu::FeOperation2D::computeGradientYTri3(cell_lid, cn_cv, in_node_coord);
63
64     Real3 node_vector_integral = area * dxU[node_lid] * dxU + area * dyU[node_lid] * dyU;
65     return {node_vector_integral[0], node_vector_integral[1], node_vector_integral[2]};
66 }
67
```

# Backup1

```
14 /*-----*/
15 /**
16 * @brief Computes the element matrix for a triangular element (P1 FE).
17 *
18 * This function calculates the integral of the expression:
19 *  $\int \int (\frac{\partial u}{\partial x} \frac{\partial v}{\partial x} + \frac{\partial u}{\partial y} \frac{\partial v}{\partial y}) d\Omega$ 
20 *
21 * Steps involved:
22 * 1. Calculate the area of the triangle.
23 * 2. Compute the gradients of the shape functions.
24 * 3. Return  $\int \int (\frac{\partial u}{\partial x} \frac{\partial v}{\partial x} + \frac{\partial u}{\partial y} \frac{\partial v}{\partial y}) d\Omega$ .
25 */
26 /*-----*/
27
28 RealMatrix<3, 3> FemModule::computeElementMatrixTria3(Cell cell)
29 {
30     Real area = ArcaneFemFunctions::MeshOperation::computeAreaTria3(cell, m_node_coord);
31
32     Real3 dxU = ArcaneFemFunctions::FeOperation2D::computeGradientXTria3(cell, m_node_coord);
33     Real3 dyU = ArcaneFemFunctions::FeOperation2D::computeGradientYTri3(cell, m_node_coord);
34
35     return area * (dxU ^ dxU) + area * (dyU ^ dyU);
36 }
37
38 /*-----*/
39 /**
40 * @brief ...
41 */
42 ARCCORE_HOST_DEVICE RealMatrix<3, 3>
43 _computeElementMatrixTria3Gpu(CellLocalId cell_lid, const IndexedCellNodeConnectivityView& cn_cv, const ax::VariableNodeReal3InView& in_node_coord);
44 {
45     Real area = Arcane::FemUtils::Gpu::MeshOperation::computeAreaTria3(cell_lid, cn_cv, in_node_coord);
46
47     Real3 dxU = FemUtils::Gpu::FeOperation2D::computeGradientXTria3(cell_lid, cn_cv, in_node_coord);
48     Real3 dyU = FemUtils::Gpu::FeOperation2D::computeGradientYTri3(cell_lid, cn_cv, in_node_coord);
49
50     return area * (dxU ^ dxU) + area * (dyU ^ dyU);
51 }
52
53 /*-----*/
54 /**
55 */
56 ARCCORE_HOST_DEVICE RealMatrix<1, 3>
57 _computeElementVectorTria3Gpu(CellLocalId cell_lid, const IndexedCellNodeConnectivityView& cn_cv, const ax::VariableNodeReal3InView& in_node_coord, Int32 node_lid)
58 {
59     Real area = Arcane::FemUtils::Gpu::MeshOperation::computeAreaTria3(cell_lid, cn_cv, in_node_coord);
60
61     Real3 dxU = FemUtils::Gpu::FeOperation2D::computeGradientXTria3(cell_lid, cn_cv, in_node_coord);
62     Real3 dyU = FemUtils::Gpu::FeOperation2D::computeGradientYTri3(cell_lid, cn_cv, in_node_coord);
63
64     Real3 node_vector_integral = area * dxU[node_lid] * dxU + area * dyU[node_lid] * dyU;
65     return {node_vector_integral[0], node_vector_integral[1], node_vector_integral[2]};
66 }
67
```

CPU

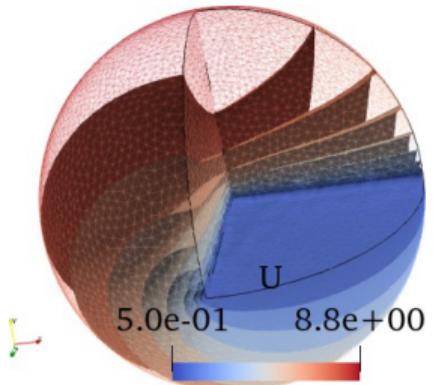
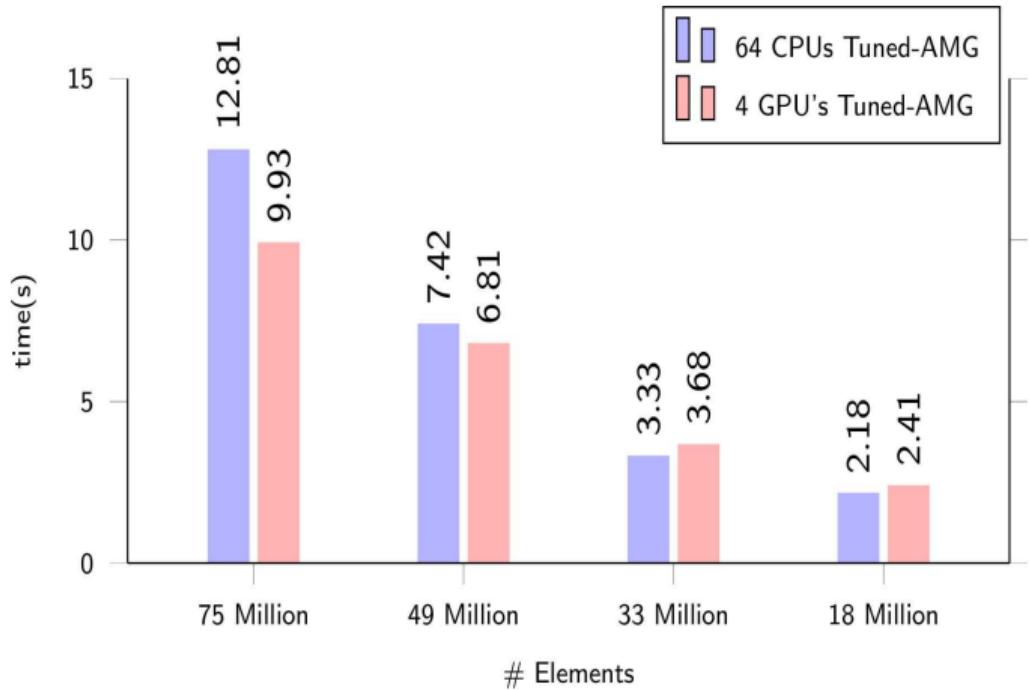
GPU

AF-GPU



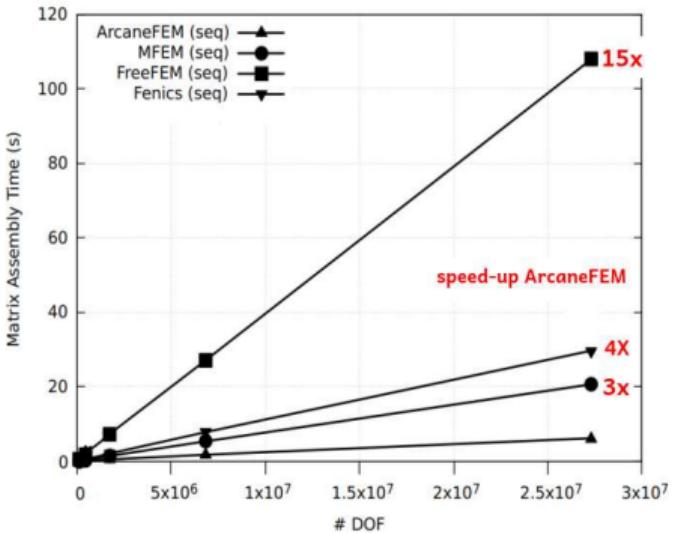


## Backup2: Large problems are needed for GPU

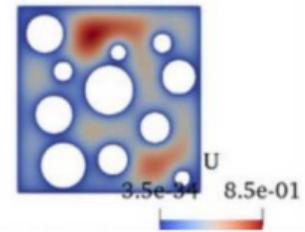




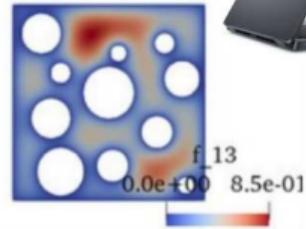
## Backup3: Comparison with state of the art FEM solver



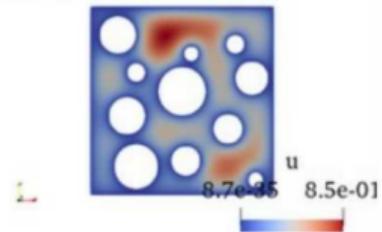
ArcaneFEM



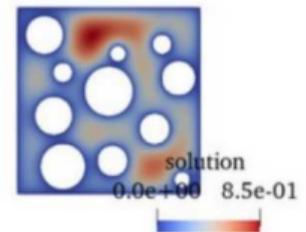
Fenics



FreeFEM

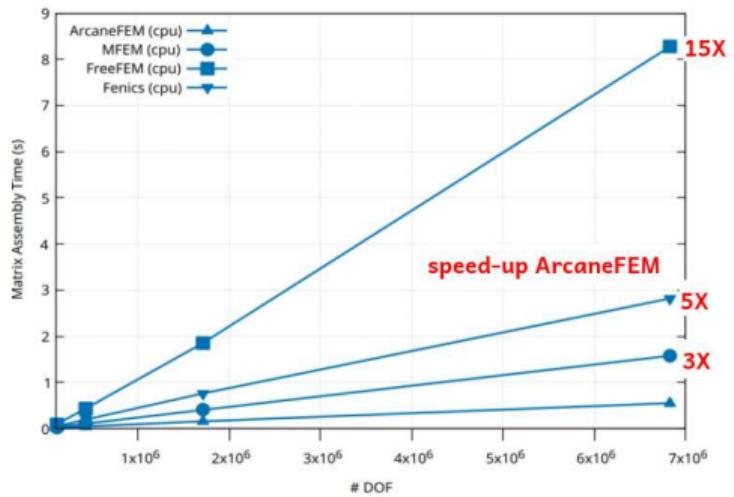


MFEM





## Backup3: Comparison with state of the art FEM solver

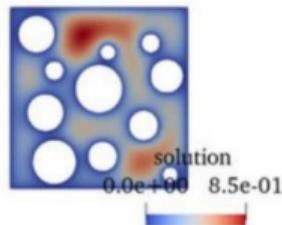
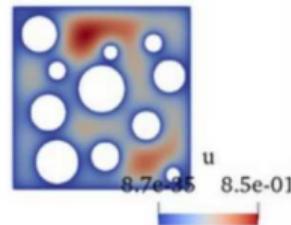
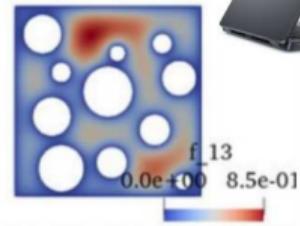
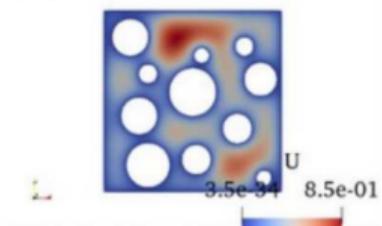


ArcaneFEM

FreeFEM

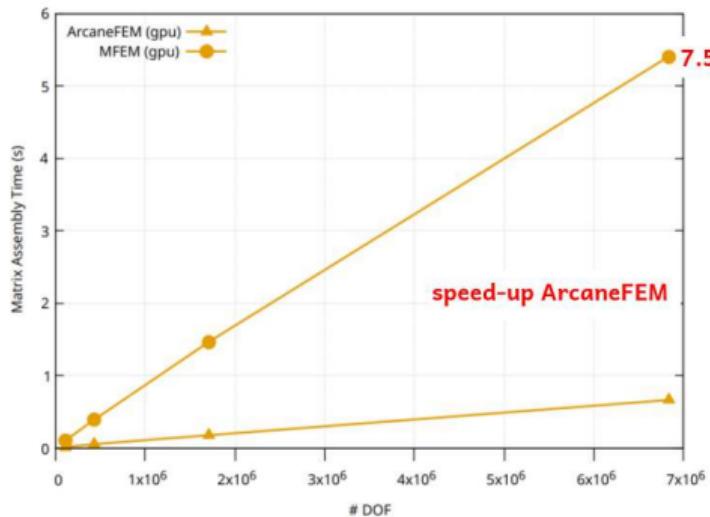
Fenics

MFEM

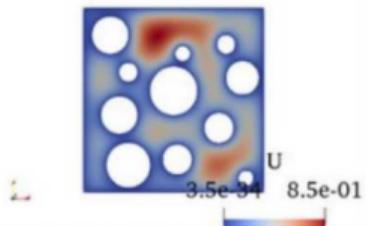




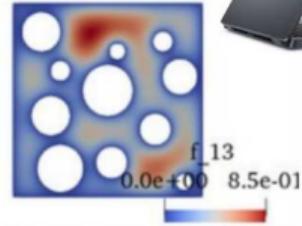
## Backup3: Comparison with state of the art FEM solver



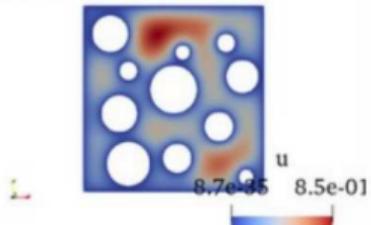
ArcaneFEM



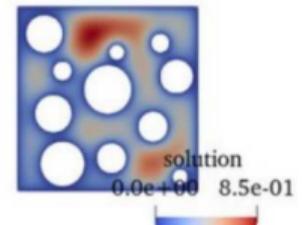
Fenics



FreeFEM

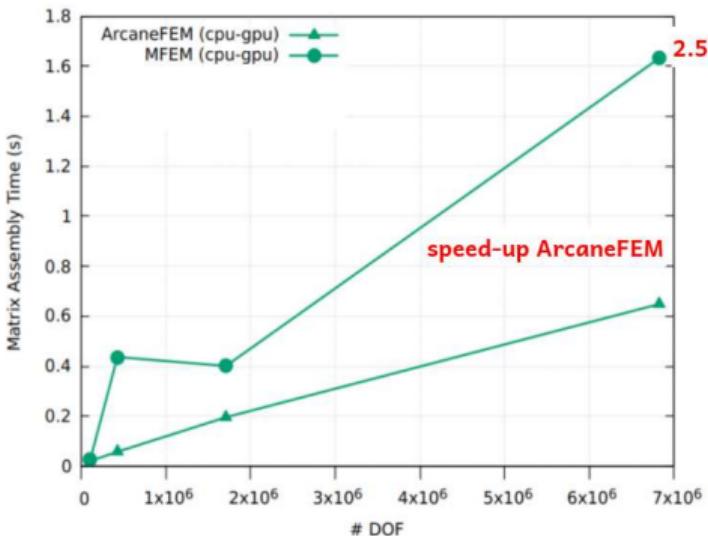


MFEM

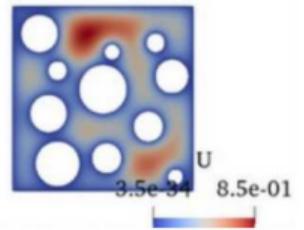




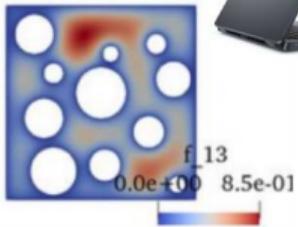
## Backup3: Comparison with state of the art FEM solver



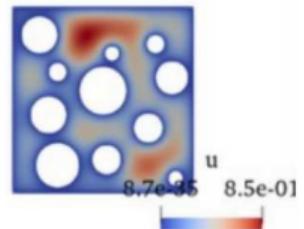
ArcaneFEM



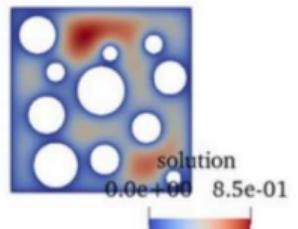
Fenics



FreeFEM

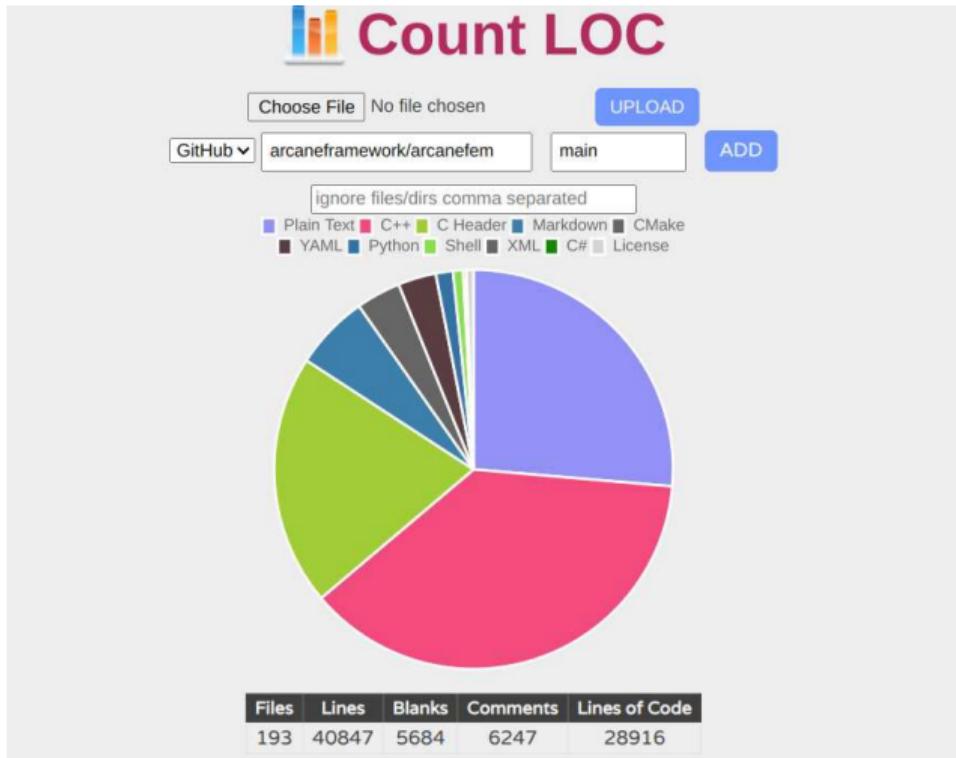


MFEM



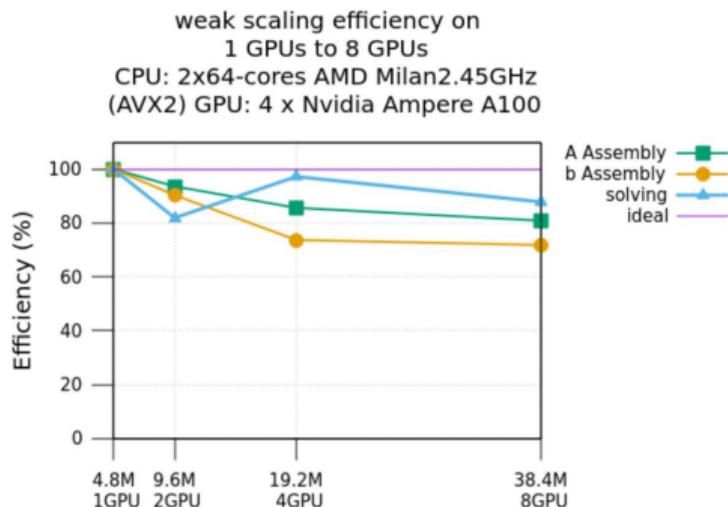
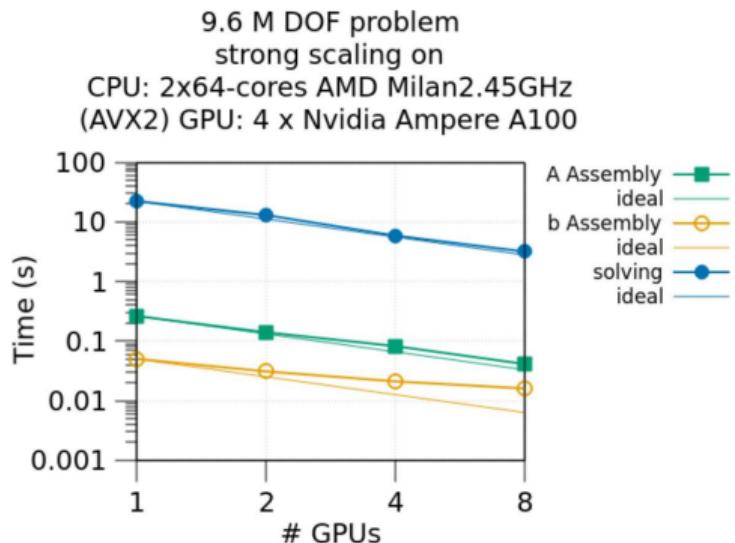


## Backup3: Lines of code





## Backup4: Scaling weak and strong NVIDIA





## Backup5: Elasticity vs Poisson

